

Data Structures – Week #6

Special Trees

Outline

- Adelson-Velskii-Landis (AVL) Trees
- Splay Trees
- B-Trees

AVL Trees

March 10, 2021

Borahan Tümer, Ph.D.

3

Motivation for AVL Trees

- Accessing a node in a BST takes $O(\log_2 n)$ in average.
- A BST can be structured so as to have an average access time of $O(n)$. *Can you think of one such BST?*
- Q: Is there a way to *guarantee a worst-case access time of $O(\log_2 n)$ per node* or can we find a way to *guarantee a BST depth of $O(\log_2 n)$* ?
- A: **AVL Trees**

March 10, 2021

Borahan Tümer, Ph.D.

4

Definition

An *AVL tree* is a *BST* with the following *balance condition*:
for each node in the BST, the height of left and right sub-trees can differ by at most 1, or

$$|h_{N_L} - h_{N_R}| \leq 1.$$

March 10, 2021

Borahan Tümer, Ph.D.

5

Remarks on Balance Condition

- *Balance condition must be easy to maintain*:
 - This is the reason, for example, for the balance condition's not being as follows: the height of left and right sub-trees of each node have the same height.
- *It ensures the depth of the BST is $O(\log_2 n)$.*
- The *height information is stored* as an additional field in `BTNodeType`.

March 10, 2021

Borahan Tümer, Ph.D.

6

Structure of an AVL Tree

```
struct BTreeNodeType {  
    infoType *data;  
    unsigned int height;  
    struct BTreeNodeType *left;  
    struct BTreeNodeType *right;  
}
```

March 10, 2021

Borahan Tümer, Ph.D.

7

Rotations

Definition:

- *Rotation* is the operation performed on a BST to restore its AVL property lost as a result of an insert operation.
- We consider the node α whose new balance violates the AVL condition.

March 10, 2021

Borahan Tümer, Ph.D.

8

Rotation


- Violation of AVL condition
- The AVL condition violation may occur in four cases:
 - Insertion into *left subtree of the left child* (L/L)
 - Insertion into *right subtree of the left child* (R/L)
 - Insertion into *left subtree of the right child* (L/R)
 - Insertion into *right subtree of the right child* (R/R)
- The outside cases 1 and 4 (i.e., L/L and R/R) are fixed by a *single rotation*.
- The other cases (i.e., R/L and L/R) need two rotations called *double rotation* to get fixed.
- These are fundamental operations in balanced-tree algorithms.

March 10, 2021

Borahan Tümer, Ph.D.

9

Single Rotation (L/L)




March 10, 2021

Borahan Tümer, Ph.D.

10

Single Rotation (R/R)




March 10, 2021

Borahan Tümer, Ph.D.

11

Double Rotation (R/L)




Single rotation cannot fix the AVL condition violation!!!

March 10, 2021

Borahan Tümer, Ph.D.

12

Double Rotation (R/L)



The symmetric case (L/R) is handled similarly left as an exercise to you!

March 10, 2021

Borahan Tümer, Ph.D.

13

Constructing an AVL Tree – Animation

48

48


March 10, 2021

Borahan Tümer, Ph.D.

14

Constructing an AVL Tree – Animation

48 16




March 10, 2021

Borahan Tümer, Ph.D.

15

Constructing an AVL Tree – Animation

48 16 24




March 10, 2021

Borahan Tümer, Ph.D.

16

Constructing an AVL Tree – Animation

48 16 24




March 10, 2021

Borahan Tümer, Ph.D.

17

Constructing an AVL Tree – Animation

48 16 24 20




March 10, 2021

Borahan Tümer, Ph.D.

18

Constructing an AVL Tree – Animation

48 16 24 20 8




March 10, 2021

Borahan Tümer, Ph.D.

19

Constructing an AVL Tree – Animation

48 16 24 20 8 12




March 10, 2021

Borahan Tümer, Ph.D.

20

Constructing an AVL Tree – Animation

48 16 24 20 8 12




March 10, 2021

Borahan Tümer, Ph.D.

21

Constructing an AVL Tree – Animation

48 16 24 20 8 12 32




March 10, 2021

Borahan Tümer, Ph.D.

22

Constructing an AVL Tree – Animation

48 16 24 20 8 12 32 54




March 10, 2021

Borahan Tümer, Ph.D.

23

Constructing an AVL Tree – Animation

48 16 24 20 8 12 32 54 72




March 10, 2021

Borahan Tümer, Ph.D.

24

Constructing an AVL Tree – Animation

48 16 24 20 8 12 32 54 72




March 10, 2021

Borahan Tümer, Ph.D.

25

Constructing an AVL Tree – Animation

48 16 24 20 8 12 32 54 72 18




March 10, 2021

Borahan Tümer, Ph.D.

26

Constructing an AVL Tree – Animation

48 16 24 20 8 12 32 54 72 18




March 10, 2021

Borahan Tümer, Ph.D.

27

Constructing an AVL Tree – Animation

48 16 24 20 8 12 32 54 72 18 96




March 10, 2021

Borahan Tümer, Ph.D.

28

Constructing an AVL Tree – Animation

48 16 24 20 8 12 32 54 72 18 96




March 10, 2021

Borahan Tümer, Ph.D.

29

Constructing an AVL Tree – Animation

48 16 24 20 8 12 32 54 72 18 96 64 17 60 98 68 84 36 30



March 10, 2021

Borahan Tümer, Ph.D.

30

Height versus Number of Nodes

- The *minimum number* of nodes in an AVL tree recursively relates to the height of the tree as follows:

$$S(h) = S(h-1) + S(h-2) + 1;$$

$$\text{Initial Values: } S(0)=1; S(1)=2$$

Homework: Solve for $S(h)$ as a function of h !

March 10, 2021

Borahan Tümer, Ph.D.

31

Splay Trees

March 10, 2021

Borahan Tümer, Ph.D.

32

Motivation for Splay Trees

- We are looking for a data structure where, *even though some worst case ($O(n)$) accesses may be possible, m consecutive tree operations starting from an empty tree (inserts, finds and/or removals) take $O(m \cdot \log_2 n)$.*
- Here, the main idea is to assume that, *$O(n)$ accesses are not bad as long as they occur relatively infrequently.*
- Hence, we are looking for *modifications of a BST per tree operation that attempts to minimize $O(n)$ accesses.*

March 10, 2021

Borahan Tümer, Ph.D.

33

Splaying

- The underlying idea of splaying is to *move a deep node accessed upwards to the root*, assuming that it will be accessed in the near future again.
- While doing this, other deep nodes are also carried up to smaller depth levels, making the average depth of nodes closer to *$O(\log_2 n)$.*

March 10, 2021

Borahan Tümer, Ph.D.

34

Splaying


- Splaying is similar to bottom-up AVL rotations
 - If a node *X is the child of the root R*,
 - then we *rotate only X and R, and this is the last rotation performed.*
- else consider *X*, its *parent P* and *grandparent G*.
Two cases and their symmetries to consider
Zig-zag case, and
Zig-zig case.

March 10, 2021

Borahan Tümer, Ph.D.

35

Zig-zag case




This is the same operation as an AVL double rotation in an R/L violation.

March 10, 2021

Borahan Tümer, Ph.D.

36

Zig-zig case




March 10, 2021

Borahan Tümer, Ph.D.

37

Animated Example

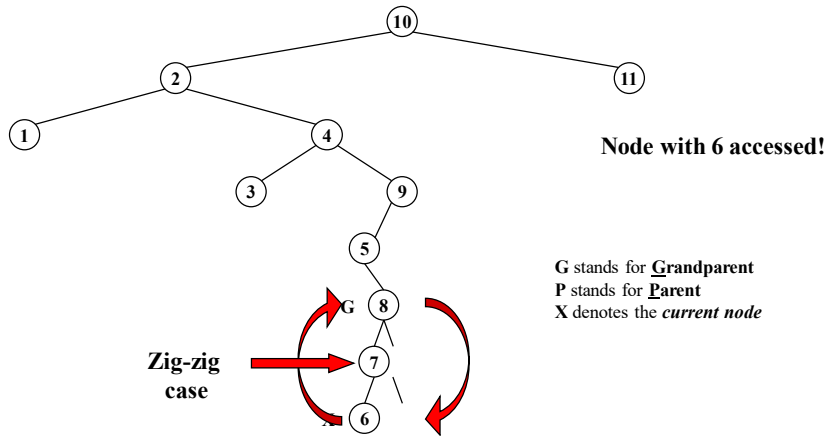


March 10, 2021

Borahan Tümer, Ph.D.

38

Animated Example

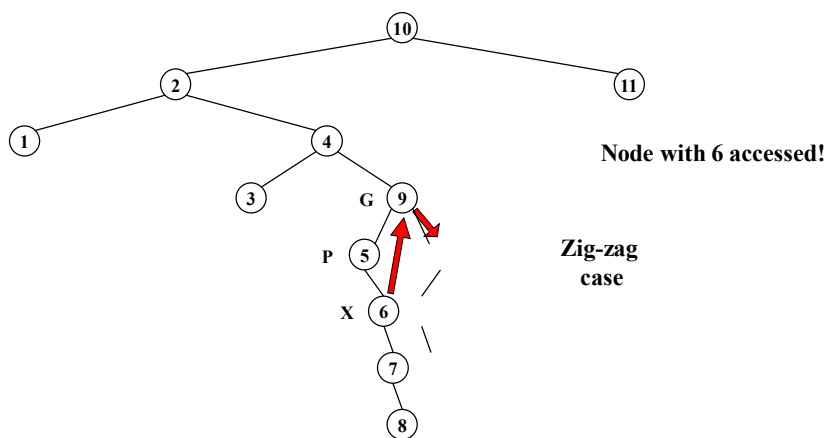


March 10, 2021

Borahan Tümer, Ph.D.

39

Animated Example




March 10, 2021

Borahan Tümer, Ph.D.

40

Animated Example




March 10, 2021

Borahan Tümer, Ph.D.

41

Animated Example




March 10, 2021

Borahan Tümer, Ph.D.

42

Animated Example



Node with 6 accessed!

March 10, 2021

Borahan Tümer, Ph.D.

43

B-Trees

March 10, 2021

Borahan Tümer, Ph.D.

44

Motivation for B-Trees

- Two technologies for providing memory capacity in a computer system
 - *Primary (main) memory (silicon chips)*
 - *Secondary storage (magnetic disks)*
- Primary memory
 - 5 orders of magnitude (i.e., about 10^5 times) *faster*,
 - 2 orders of magnitude (about 100 times) *more expensive*, and
 - by at least 2 orders of magnitude *less in size*

than secondary storage due to mechanical operations involved in magnetic disks.

March 10, 2021

Borahan Tümer, Ph.D.

45

Motivation for B-Trees

- During one disk read or disk write (4-8.5msec for 7200 RPM disks), MM can be accessed about 10^5 times (100 nanosec per access).
- To reimburse (compensate) for this time, at each disks access, *not a single item*, but one or more *equal-sized pages* of items (each page 2^{11} - 2^{14} bytes) are accessed.
- We need some data structure to store these *equal sized pages* in MM.
- *B-Trees*, with their *equal-sized leaves (as big as a page)*, are suitable data structures for storing and performing regular operations on paged data.

March 10, 2021

Borahan Tümer, Ph.D.

46

B-Trees

- A *B-tree* is a rooted tree with the following properties:
- Every node x has the following fields:
 - $n[x]$, the number of keys currently stored in x .
 - the $n[x]$ keys themselves, in *non-decreasing order*, so that

$$key_1[x] \leq key_2[x] \leq \dots \leq key_{n[x]}[x],$$
 - $leaf[x]$, a boolean value, true if x is a leaf.

March 10, 2021

Borahan Tümer, Ph.D.

47

B-Trees

- Each internal node has $n[x]+1$ pointers, $c_1[x], \dots, c_{n[x]+1}[x]$, to its children. *Leaf nodes* have *no children*, hence no pointers!
- The keys separate the ranges of keys stored in each subtree: if k_i is any key stored in the subtree with root $c_i[x]$, then

$$k_i \leq key_1[x] \leq key_2[x] \leq \dots \leq key_{n[x]}[x] \leq k_{n[x]+1}.$$
- *All leaves have the same depth, h , equal to the tree's height.*

March 10, 2021

Borahan Tümer, Ph.D.

48

B-Trees

- There are lower and upper bounds on the number of keys a node may contain. These bounds can be expressed in terms of a fixed integer $t \geq 2$ called the *minimum degree* of the B-Tree.
 - Lower limits
 - All *nodes but the root* has *at least $t-1$* keys.
 - Every *internal node but the root* has *at least t children*.
 - A non-empty tree's **root** must have *at least one key*.

March 10, 2021

Borahan Tümer, Ph.D.

49

B-Trees

- Upper limits
 - Every *node* can contain *at most $2t-1$ keys*.
 - Every *internal node* can have *at most $2t$ children*.
 - A node is defined to be full if it has exactly *$2t-1$ keys*.
- For a *B-tree* of minimum degree $t \geq 2$ and n nodes

$$h \leq \log_t \frac{n+1}{2}$$

March 10, 2021

Borahan Tümer, Ph.D.

50

Basic Operations on B-Trees

- *B-tree search*
- *B-tree insert*
- *B-tree removal*

March 10, 2021

Borahan Tümer, Ph.D.

51

Disk Operations in B-Tree operations

- Suppose x is a pointer to an object.
- It is accessible if it is in the main memory.
- If it is on the disk, it needs to be transferred to the main memory to be accessible. This is done by *DISK_READ(x)*.
- To save any changes made to any field(s) of the object pointed to by x , a *DISK_WRITE(x)* operation is performed.

March 10, 2021

Borahan Tümer, Ph.D.

52




Figure 4.63 B-tree of order 5




Figure 4.64 B-tree after insertion of 57 into the tree in Figure 4.63

Marc

53




Figure 4.65 Insertion of 55 into the B-tree in Figure 4.64 causes a split into two leaves




Figure 4.66 Insertion of 40 into the B-tree in Figure 4.65 causes a split into two leaves and then a split of the parent node

54




Figure 4.66 Insertion of 40 into the B-tree in Figure 4.65 causes a split into two leaves and then a split of the parent node




Figure 4.67 B-tree after the deletion of 99 from the B-tree in Figure 4.66


March 10, 2021

Borahan Tümer, Ph.D.

55

Search in B-Trees

- Similar to search in BSTs with the exception that instead of a binary, a multi-way ($n[x]+1$ -way) decision is made.



March 10, 2021

Borahan Tümer, Ph.D.

56

Search in B-Trees

```

B-tree-Search(x, k)
{
  i=1;
  while (i ≤ n[x] and k > keyi[x]) i++;
  if (i ≤ n[x] and k = keyi[x])           // if key found
    return (x,i);
  if (leaf[x])                           // if key not found at a leaf
    return NULL;
  else {DISK_READ(ci[x]);                // if key < keyi[x]
        return B-tree-Search(ci[x],k);}
}

```

March 10, 2021

Borahan Tümer, Ph.D.

57

Insertion in B-Trees

- Insertion into a B-tree is more complicated than that into a BST, since the creation of a new node to place the new key may violate the B-tree property of the tree.
- Instead, the key is put *into a leaf node x if it is not full*.
- If full, a *split* is performed, which splits a full node (with *2t-1* keys) at its *median key*, *key_t[x]*, into two nodes with *t-1* keys each.
- *key_t[x]* moves up into the parent of *x* and identifies the split point of the two new trees.

March 10, 2021

Borahan Tümer, Ph.D.

58

Insertion in B-Trees

- A *single-pass insertion* starts at the root traversing *down to the leaf* into which the key is to be inserted.
- On the path down, *all full nodes are split* including a full leaf that also guarantees a parent with an available position for the median key of a full node to be placed.


March 10, 2021

Borahan Tümer, Ph.D.

59

Insertion in B-Trees: Example

69 inserted...




March 10, 2021

Borahan Tümer, Ph.D.

60


Insertion in B-Trees: Example

69 inserted...




Insertion in B-Trees: Example

69 inserted...



Insertion in B-Trees: Example

69 inserted...



March 10, 2021

Borahan Tümer, Ph.D.

67 69 70

63

Insertion in B-Trees: B-tree-Insert

B-tree-Insert(T, k)

```
{ r=root[T];
  if ( $n[r] == 2t-1$ ) {
    s=malloc(new-B-tree-node);
    root[T]=s;
    leaf[s]=false;
    n[s]=0;
     $c_1[s]=r$ ;
    B-tree-Split-Child(s,1,r);
    B-tree-Insert-Nonfull(s,k); }
  else B-tree-Insert-Nonfull(r,k);
}
```

March 10, 2021

Borahan Tümer, Ph.D.

64

65

66

Insertion in B-Trees: B-tree-Insert-Nonfull

```

B-tree-Insert-Nonfull(x,k)
{
  i=n[x];
  if (leaf[x]) {
    while (i≥1 and k < keyi[x]) {keyi+1[x]=keyi[x]; i--;}
    keyi+1[x]=k;
    n[x]++;
    DISK_WRITE(x);
  }
  else {
    while (i≥1 and k < keyi[x]) i--;
    i++;
    DISK_READ(ci[x]);
    if (n[ci[x]]==2t-1) {
      B-tree-Split-Child(x,i, ci[x]);
      if (k > keyi[x]) i++;
    }
    B-tree-Insert-Nonfull(ci[x],k);
  }
}

```

if x is a leaf
 then place key in x;
 write x on disk;
 else find the node (root of subtree) key goes to;
 read node from disk;
 if node full
 split node at key's position;
 recursive call with node split and key;

March 10, 2021

Borahan Tümer, Ph.D.

69

Removing a key from a B-Tree

- Removal in B-trees is different than insertion only in that *a key may be removed from any node, not just from a leaf.*
- As the insertion algorithm splits any full node down the path to the leaf to which the key is to be inserted, a recursive removal algorithm may be written to ensure that for any call to removal on a node x , the number of keys in x is at least the minimum degree t .

March 10, 2021

Borahan Tümer, Ph.D.

70

Various Cases of Removing a key from a B-Tree

1. If the key k is in node x and x is a leaf, remove the key k from x .
2. If the key k is in node x and x is an internal node, then
 - a. If the child y that precedes k in node x has at least t keys, then find the predecessor k' of k in the subtree rooted at y . Recursively delete k' , and replace k by k' in x . Finding k' and deleting it can be performed in a single downward pass.

March 10, 2021

Borahan Tümer, Ph.D.

71

Various Cases of Removal a key from a B-Tree

- b. Symmetrically, if the child z that follows k in node x has at least t keys, then find the successor k' of k in the subtree rooted at z . Recursively delete k' , and replace k by k' in x . Finding k' and deleting it can be performed in a single downward pass.
- c. Otherwise, if both y and z have only $t-1$ keys, merge k and all of z into y so that x loses both k and the pointer to z and y now contains $2t-1$ keys. Free z and recursively delete k from y .

March 10, 2021

Borahan Tümer, Ph.D.

72

Various Cases of Removal a key from a B-Tree

3. If k is not present in internal node x , determine root $c_i[x]$ of the subtree that must contain k , if k exists in the tree. If $c_i[x]$ has only $t-1$ keys, execute step 3a or 3b as necessary to guarantee that we descend to a node containing at least t keys. Then finish by recursing on the appropriate child of x .

March 10, 2021

Borahan Tümer, Ph.D.

73

Various Cases of Removal a key from a B-Tree

- a. If $c_i[x]$ has only $t-1$ keys but has an **immediate sibling with at least t keys**, give $c_i[x]$ an extra key by moving a key **from x down into $c_i[x]$** , moving a key **from $c_i[x]$'s immediate left or right sibling up into x** , and moving the appropriate child pointer from the sibling into $c_i[x]$.
- b. If $c_i[x]$ and both of $c_i[x]$'s immediate siblings have $t-1$ keys, merge $c_i[x]$ with one sibling, which involves moving a key from x down into the new merged node to become the *median key* for that node.

March 10, 2021

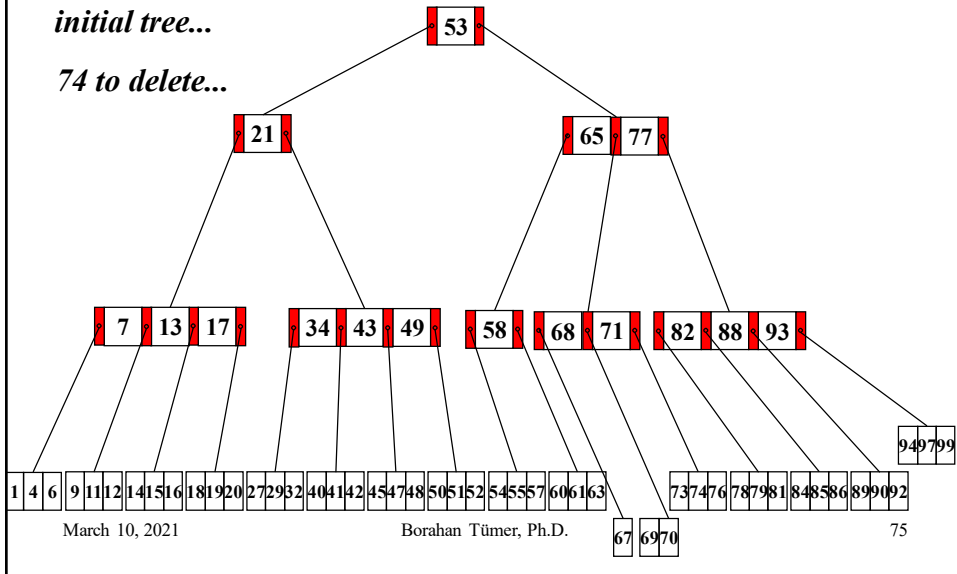
Borahan Tümer, Ph.D.

74

Removal in B-Trees: Example

initial tree...


74 to delete...



Removal in B-Trees: Example


74 removed...

21 to delete...



Removal in B-Trees: Example


21 removed...
65 to remove...



Removal in B-Trees: Example

65 removed...
69 to remove...


Case 2b recursively followed by case 3a



Removal in B-Trees: Example

69 removed...

Case 2c recursively
followed by case 1



March 10, 2021

Borahan Tümer, Ph.D.

79