

# Course Overview

**CSE 238/2038/2138: Systems Programming**

**Instructor:**

Fatma CORUT ERGİN

*Slides adapted from Bryant & O'Hallaron's slides*

# Overview

- Course logistics
- Course outcomes
- Course coverage

# Fatma CORUT ERGİN

- e-mail: [fatma.ergin@marmara.edu.tr](mailto:fatma.ergin@marmara.edu.tr)
- Office: MB 449
- Zoom link: <https://cutt.ly/4ILV0ci>
  - Meeting ID: 255 133 3127
  - Passcode: 238495

# Teaching Assistants

- Zuhall ALTUNTAŞ
  - [zuhal.altuntas@marmara.edu.tr](mailto:zuhal.altuntas@marmara.edu.tr)
- Lokman ALTIN
  - [lokman.altin@marmara.edu.tr](mailto:lokman.altin@marmara.edu.tr)

# Logistics

## ■ Lecture Hours

- Mondays 14.00 – 15.50 (2 sessions)
- Tuesdays 13.00 – 13.50 (1 session)

## ■ Course Content

- Canvas <https://canvas.instructure.com/enroll/CWRL68>
- You can enroll in the course with the link
- [ues.marmara.edu.tr](https://ues.marmara.edu.tr)

## ■ Lecture Videos

- [ues.marmara.edu.tr](https://ues.marmara.edu.tr)

## ■ Announcements

- Canvas <https://canvas.instructure.com/courses/2609650>
- [ues.marmara.edu.tr](https://ues.marmara.edu.tr)

# Textbook

- Randal E. Bryant and David R. O'Hallaron,
  - *Computer Systems: A Programmer's Perspective*, **Third Edition** (CS:APP3e), Pearson, 2016
  - <http://csapp.cs.cmu.edu>
- **This book really matters for the course!**
  - How to solve labs
  - Practice problems typical of exam problems

# Grading (tentative)

- 70% attendance is a must to pass the course
- Scheduled quizzes – each 2% (on Canvas)
  - You may have scheduled quizzes.
  - You will take quizzes on Canvas
  - The quiz date will be announced at least 1 day before.
  - These quizzes will take between 10-20 minutes
- Pop quizzes – 10% (on Zoom lecture)
  - Pop quizzes can happen any time during the lecture hours.
  - The questions will be asked on Zoom lecture meeting.
  - They will be on the topic of the corresponding lecture.
  - These quizzes will be very-short (1-2 minutes)
- 3 Projects – 30% (may change according to scheduled quiz number)
- Midterm Exam – 20% (on Zoom)
- Final Exam – 40% (on Zoom)

# Canvas Participation

- You will receive an e-mail from “Instructure Canvas” as an invitation
  - You should accept the invitation
- Your scheduled quizzes will be prepared on Canvas, so it is very important for you to enroll in the class as soon as possible



# Cheating: Description

## ■ What is cheating?

- Sharing code: by copying, retyping, looking at, or supplying a file
- Describing: verbal description of code from one person to another.
- Coaching: helping your friend to write a code, line by line
- Searching the Web for solutions
- Copying code from any kind of source
  - You are only allowed to use code we supply

## ■ What is NOT cheating?

- Explaining how to use systems or tools
- Helping others with high-level design issues

# Cheating: Consequences

- Penalty for cheating:
  - “0” grade for both parties
- Detection of cheating:
  - We have sophisticated tools for detecting code plagiarism

**Start early**

**Ask the staff for help when you get stuck**

Welcome and Enjoy!

# Course Theme:

## Systems Knowledge is Power!

### ■ Systems knowledge

- How hardware (processors, memories, disk drives, network infrastructure) plus software (operating systems, compilers, libraries, network protocols) combine to support the execution of application programs
- How you (as a programmer) can best use these resources

### ■ Useful outcomes of the course

- Become more effective programmers
  - Able to find and eliminate bugs efficiently
  - Able to understand and tune for program performance
- Prepare for later “systems” classes in CSE
  - Computer Organization, Operating Systems, etc.

# Understand How Things Work

- Why do I need to know this stuff?
  - Abstraction is good, but don't forget reality
- Most CSE courses emphasize abstraction
  - Abstract data types
  - Asymptotic analysis
- These abstractions have limits
  - Especially in the presence of bugs
  - Need to understand details of underlying implementations
  - Sometimes the abstract interfaces don't provide the level of control or performance you need

# Great Reality #1:

## Ints are not Integers, Floats are not Reals

### ■ Example 1: Is $x^2 \geq 0$ ?

- Float's: Yes!

- Int's:

- $40000 * 40000$   
= 1600000000
- $50000 * 50000$   
= ??



### ■ Example 2: Is $(x + y) + z = x + (y + z)$ ?

- Unsigned & Signed Int's: Yes!

- Float's:

- $(1e20 + -1e20) + 3.14 \rightarrow 3.14$
- $1e20 + (-1e20 + 3.14) \rightarrow ??$

# Computer Arithmetic

## ■ Does not generate random values

- Arithmetic operations have important mathematical properties

## ■ Cannot assume all “usual” mathematical properties

- Due to finiteness of representations
- Integer operations satisfy “ring” properties
  - Commutativity, associativity, distributivity
- Floating point operations satisfy “ordering” properties
  - Monotonicity, values of signs

## ■ Observation

- Need to understand which abstractions apply in which contexts
- Important issues for compiler writers and serious application programmers

# Great Reality #2:

## You've Got to Know Assembly

- **You'll never write programs in assembly**
  - Compilers are much better & more patient than you are
- **But: Understanding assembly is key to machine-level execution model**
  - Behavior of programs in presence of bugs
    - High-level language models break down
  - Tuning program performance
    - Understand optimizations done / not done by the compiler
    - Understanding sources of program inefficiency
  - Implementing system software
    - Compiler has machine code as target
    - Operating systems must manage process state
  - Creating / fighting malware
    - x86 assembly is the language of choice!



# Great Reality #3: Memory Matters

## Random Access Memory Is an Unphysical Abstraction

### ■ Memory is not unbounded

- It must be allocated and managed
- Many applications are memory dominated

### ■ Memory referencing bugs especially evil

- Effects are distant in both time and space

### ■ Memory performance is not uniform

- Cache and virtual memory effects can greatly affect program performance
- Adapting program to characteristics of memory system can lead to major speed improvements

# Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;  
  
double fun(int i) {  
    struct_t s;  
    s.d = 3.14;  
    s.a[i] = 1073741121;  
    return s.d;  
}
```

fun(0) →	3.14	(a[0]=1073741121, a[1]=0)
fun(1) →	3.14	(a[0]=0, a[1]=1073741121)
fun(2) →	3.139999866485284	(a[0]=0, a[1]=0)
fun(3) →	1.999329872131348	(a[0]=0, a[1]=0)
fun(4) →	3.14	(a[0]=0, a[1]=0)
fun(5) →	3.14	(a[0]=0, a[1]=0)
fun(6) →	3.14	(a[0]=0, a[1]=0)
fun(7) →	Segmentation fault	

■ Result is system specific

# Memory Referencing Bug Example

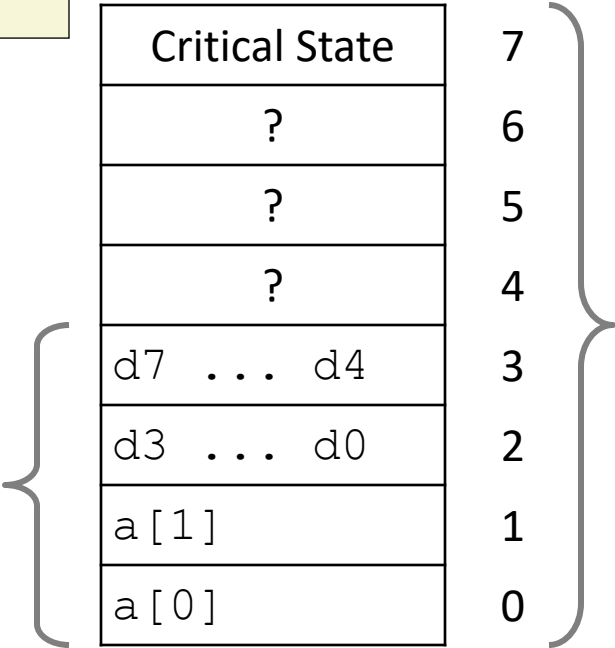
```
typedef struct {
    int a[2];
    double d;
} struct_t;

double fun(int i) {
    struct_t s;
    s.d = 3.14;
    s.a[i] = 1073741121;
    return s.d;
}
```

- fun(0) → 3.14
- fun(1) → 3.14
- fun(2) → 3.1399999866485284
- fun(3) → 1.999329872131348
- fun(4) → 3.14
- fun(5) → 3.14
- fun(6) → 3.14
- fun(7) → Segmentation fault

Explanation:

struct\_t



Location accessed by fun(i)

# What About This?

```
typedef struct {  
    double d;  
    int a[2];  
} struct_t  
  
double fun(int i) {  
    struct_t s;  
    s.d = 3.14;  
    s.a[i] = 1073741121;  
    return s.d;  
}
```

fun(0)	→	3.14
fun(1)	→	3.14
fun(2)	→	3.14
fun(3)	→	3.14
fun(4)	→	3.14
fun(5)	→	Segmentation fault

Explanation:

struct\_t

Critical State	7
?	6
?	5
?	4
a[1]	3
a[0]	2
d7 ... d4	1
d3 ... d0	0

Location accessed by  
fun(i)

# Memory Referencing Errors

## ■ C and C++ do not provide any memory protection

- Out of bounds array references
- Invalid pointer values
- Abuses of malloc/free

## ■ Can lead to nasty bugs

- Whether or not bug has any effect depends on system and compiler
- Action at a distance
  - Corrupted object logically unrelated to one being accessed
  - Effect of bug may be first observed long after it is generated

## ■ How can I deal with this?

- Program in Java, Ruby, Python, ML, ...
- Understand what possible interactions may occur
- Use or develop tools to detect referencing errors (e.g. Valgrind)

# Great Reality #4: There's more to performance than asymptotic complexity

- **Constant factors matter too!**
- **And even exact op count does not predict performance**
  - Easily see 10:1 performance range depending on how code written
  - Must optimize at multiple levels: algorithm, data representations, procedures, and loops
- **Must understand system to optimize performance**
  - How programs compiled and executed
  - How to measure program performance and identify bottlenecks
  - How to improve performance without destroying code modularity and generality

# Memory System Performance Example

```
void copyij(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```

4.3ms

```
void copyji(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```

81.8ms

2.0 GHz Intel Core i7 Haswell

# Course Perspective

- Our Course is Programmer-Centric
  - Purpose is to show that by knowing more about the underlying system, one can be more effective as a programmer
  - Enable you to write programs that are more reliable and efficient