

Bits, Bytes, and Integers

CSE 238/2038/2138: Systems Programming

Instructor:

Fatma CORUT ERGİN

Slides adapted from Bryant & O'Hallaron's slides

Today: Bits, Bytes, and Integers

- **Representing information as bits**
- **Bit-level manipulations**
- **Integers**
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
 - Summary
- **Representations in memory, pointers, strings**

Computer is a binary digital system

Binary (base 2) system:

- has two states: 0 and 1

- Basic unit of information is the *binary digit*, or *bit*.
- Values with more than two states require multiple bits.
 - A collection of **two** bits has **four** possible states:
00, 01, 10, 11
 - A collection of **three** bits has **eight** possible states:
000, 001, 010, 011, 100, 101, 110, 111
 - A collection of **n** bits has **2^n** possible states.

What kinds of data do we need to represent?

■ Everything is bits

- **Numbers** – signed, unsigned, integers, floating point, complex, rational, irrational, ...
- **Text** – characters, strings, ...
- **Images** – pixels, colors, shapes, ...
- **Sound**
- **Logical** – true, false
- **Instructions**
- ...

■ Base 2 Number Representation

- Represent 15213_{10} as 11101101101101_2
- Represent 1.20_{10} as $1.0011001100110011[0011]..._2$
- Represent 1.5213×10^4 as $1.1101101101101_2 \times 2^4$

Hexadecimal Notation

- It is often convenient

to write binary (base-2) numbers
as hexadecimal (base-16) numbers instead.

- fewer digits -- four bits per hex digit
- less error prone -- easy to corrupt long string of 1's and 0's

- C notation for hexadecimal numbers:

- `0xFA1D37B`
- `0xfa1d37b`

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

***This is not a new machine representation,
just a convenient way to write the number.***

Converting from Binary to Hexadecimal

- Every four bits is a hex digit.
 - start grouping from right-hand side

<u>011</u>	<u>1010</u>	<u>1000</u>	<u>1111</u>	<u>0100</u>	<u>1101</u>	<u>0111</u>
↓	↓	↓	↓	↓	↓	↓
3	A	8	F	4	D	7

Sizes of Data Types (in bytes)

C Data Type	Typical 32-bit	Typical 64-bit	x86-64
<code>char</code>	1	1	1
<code>short</code>	2	2	2
<code>int</code>	4	4	4
<code>long</code>	4	8	8
<code>float</code>	4	4	4
<code>double</code>	8	8	8
<code>long double</code>	–	–	10/16
<code>pointer</code>	4	8	8

1 byte = 8 bits

Today: Bits, Bytes, and Integers

- Representing information as bits
- **Bit-level manipulations**
- **Integers**
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
 - Summary
- Representations in memory, pointers, strings

Logical Operations (Boolean Algebra)

■ Operations on logical TRUE or FALSE

- two states -- takes one bit to represent: TRUE=1, FALSE=0

A	B	AND	OR	XOR	A	NOT
		A & B	A B	A ^ B		~ A
0	0	0	0	0	0	1
0	1	0	1	1	1	0
1	0	0	1	1		
1	1	1	1	0		

■ View n -bit number as a collection of n logical values

- operation applied to each bit independently

General Boolean Algebras

■ Operate on Bit Vectors

- Operations applied bitwise

01101001	01101001	01101001	
& 01010101	01010101	^ 01010101	~ 01010101
<u> </u>	<u> </u>	<u> </u>	<u> </u>
01000001	01111101	00111100	10101010

■ All of the Properties of Boolean Algebra Apply

Example: Representing & Manipulating Sets

■ Representation

- Width w bit vector represents subsets of $\{0, \dots, w-1\}$

- $a_j = 1$ if $j \in A$

- 01101001 $A = \{0, 3, 5, 6\}$

76543210

- 01010101 $A = \{0, 2, 4, 6\}$

76543210

■ Operations

- | | | | |
|-----|----------------------|----------|------------------------|
| ■ & | Intersection | 01000001 | $\{0, 6\}$ |
| ■ | Union | 01111101 | $\{0, 2, 3, 4, 5, 6\}$ |
| ■ ^ | Symmetric difference | 00111100 | $\{2, 3, 4, 5\}$ |
| ■ ~ | Complement | 10101010 | $\{1, 3, 5, 7\}$ |

Bit-Level Operations in C

■ Operations `&`, `|`, `~`, `^` Available in C

- Apply to any “integral” data type
 - `long`, `int`, `short`, `char`, `unsigned`
- View arguments as bit vectors
- Arguments applied bit-wise

■ Examples

- `~ 0x41 → 0xBE`
 - `~ 0100 00012 → 1011 11102`
- `~ 0x00 → 0xFF`
 - `~ 0000 00002 → 1111 11112`
- `0x69 & 0x55 → 0x41`
 - `0110 10012 & 0101 01012 → 0100 00012`
- `0x69 | 0x55 → 0x7D`
 - `0110 10012 | 0101 01012 → 0111 11012`

Logic Operations in C

■ Contrast to Bit-Level Logical Operators

- `&&`, `||`, `!`
 - View 0 as “False”
 - Anything nonzero as “True”
 - Always return 0 or 1

■ Examples

- `!0x41` \rightarrow `0x00`
- `!0x00` \rightarrow `0x01`
- `!!0x41` \rightarrow `0x01`

- `0x69 && 0x55` \rightarrow `0x01`
- `0x69 || 0x55` \rightarrow `0x01`
- `p && *p` (avoids null pointer access)

Logic Operations in C

■ Contrast to Bit-Level Logical Operators

- `&&`, `||`, `!`
 - View 0 as “False”
 - Anything nonzero is “True”
 - Always returns 0 or 1

■ Example

- `!0x41`
- `!0x00`
- `!!0x41`
- `0x69 && 0x55`
- `0x69 || 0x55 → 0x01`
- `p && *p` (avoids null pointer access)

Watch out for difference

`&&` vs. `&` (and `||` vs. `|`)...

one of the more common errors in C programming

Shift Operations

■ Left Shift: $x \ll y$

- Shift bit-vector x left y positions
 - Throw away extra bits on left
 - Fill with 0's on right

■ Right Shift: $x \gg y$

- Shift bit-vector x right y positions
 - Throw away extra bits on right
- Logical shift
 - Fill with 0's on left
- Arithmetic shift
 - Replicate most significant bit on left

■ Undefined Behavior

- Shift amount < 0 or \geq word size

Argument x	01100010
$\ll 3$	00010000
Log. $\gg 2$	00011000
Arith. $\gg 2$	00011000

Argument x	10100010
$\ll 3$	00010000
Log. $\gg 2$	00101000
Arith. $\gg 2$	11101000

Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- **Integers**
 - **Representation: unsigned and signed**
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
 - Summary
- Representations in memory, pointers, strings

Encoding Integers

1. Unsigned

$$B2U(X) = \sum_{i=0}^{w-1} x_i \cdot 2^i$$

- A w -bit unsigned binary number can represent 2^w unique values: from 0 to (2^w-1)
 - With 3 bits ($w=3$), we can represent the integers from 0 to 7 (2^3-1)

Binary number			Decimal Value
2^2	2^1	2^0	
0	0	0	0
0	0	1	1
0	1	0	2
0	1	1	3
1	0	0	4
1	0	1	5
1	1	0	6
1	1	1	7

Encoding Integers

1. Unsigned

```
unsigned short int x = 12345;  
unsigned short int y = 53191;
```

■ C unsigned short data type (2 bytes)

	Decimal	Binary	Hex
x	12345	00110000 00111001	30 39
y	53191	11001111 11000111	CF C7

Unsigned Example

$x = 12345: 00110000 \ 00111001$
 $y = 53191: 11001111 \ 11000111$

32768 (2^{15})	16384 (2^{14})	8192 (2^{13})	4096 (2^{12})	2048 (2^{11})	1024 (2^{10})	512 (2^9)	256 (2^8)	128 (2^7)	64 (2^6)	32 (2^5)	16 (2^4)	8 (2^3)	4 (2^2)	2 (2^1)	1 (2^0)
0	0	1	1	0	0	0	0	0	0	1	1	1	0	0	1
0	0	8192	4096	0	0	0	0	0	0	32	16	8	0	0	1
1	1	0	0	1	1	1	1	1	1	0	0	0	1	1	1
32768	16384	0	0	2048	1024	512	256	128	64	0	0	0	4	2	1

12345

53191

Numeric Ranges

■ Unsigned Values

- $UMin = 0$

000...0

- $UMax = 2^w - 1$

111...1

Encoding Integers

2. Two's Complement

$$B2T(X) = -x_{w-1} \cdot 2^{w-1} + \sum_{i=0}^{w-2} x_i \cdot 2^i$$

↑
Sign Bit

- A *w-bit* unsigned binary number can represent 2^w unique values:
from (-2^{w-1}) to $(2^{w-1} - 1)$
 - With **3** bits (**$w=3$**), we can represent the integers from -4 (**-2^2**) to 3 (**2^2-1**)
- Sign Bit
 - For two's complement, most significant bit indicates sign
 - 0 for positive
 - 1 for negative

Binary number			Decimal Value
2^2	2^1	2^0	
0	0	0	0
0	0	1	1
0	1	0	2
0	1	1	3
1	0	0	-4
1	0	1	-3
1	1	0	-2
1	1	1	-1

Encoding Integers

2. Two's Complement

```
short int x = 12345;  
short int y = -12345;
```

■ C short data type (2 bytes)

	Decimal	Binary	Hex
x	12345	00110000 00111001	30 39
y	-12345	11001111 11000111	CF C7

Two's-complement Encoding Example

x = 12345: 00110000 00111001
y = -12345: 11001111 11000111

-32768 (-2 ¹⁵)	16384 (2 ¹⁴)	8192 (2 ¹³)	4096 (2 ¹²)	2048 (2 ¹¹)	1024 (2 ¹⁰)	512 (2 ⁹)	256 (2 ⁸)	128 (2 ⁷)	64 (2 ⁶)	32 (2 ⁵)	16 (2 ⁴)	8 (2 ³)	4 (2 ²)	2 (2 ¹)	1 (2 ⁰)
0	0	1	1	0	0	0	0	0	0	1	1	1	0	0	1
0	0	8192	4096	0	0	0	0	0	0	32	16	8	0	0	1
1	1	0	0	1	1	1	1	1	1	0	0	0	1	1	1
-32768	16384	0	0	2048	1024	512	256	128	64	0	0	0	4	2	1

12345

-12345

Numeric Ranges

■ Two's Complement Values

- $TMin = -2^{w-1}$

100...0

- $TMax = 2^{w-1} - 1$

011...1

- -1

111...1

Values for $W = 16$

	Decimal	Binary	Hex
UMax	65535	11111111 11111111	FF FF
TMax	32767	01111111 11111111	7F FF
TMin	-32768	10000000 00000000	80 00
-1	-1	11111111 11111111	FF FF
0	0	00000000 00000000	00 00

Values for Different Word Sizes

	W			
	8	16	32	64
UMax	255	65,535	4,294,967,295	18,446,744,073,709,551,615
TMax	127	32,767	2,147,483,647	9,223,372,036,854,775,807
TMin	-128	-32,768	-2,147,483,648	-9,223,372,036,854,775,808

■ Observations

- $|TMin| = TMax + 1$
 - Asymmetric range
- $UMax = 2 * TMax + 1$
or $(2 * |TMin| - 1)$

■ C Programming

- `#include <limits.h>`
- Declares constants, e.g.,
 - `ULONG_MAX`
 - `LONG_MAX`
 - `LONG_MIN`
- Values platform specific

Unsigned & Signed Numeric Values

X	$B2U(X)$	$B2T(X)$
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	8	-8
1001	9	-7
1010	10	-6
1011	11	-5
1100	12	-4
1101	13	-3
1110	14	-2
1111	15	-1

■ Equivalence

- Same encodings for nonnegative values

■ Uniqueness

- Every bit pattern represents unique integer value
- Each representable integer has unique bit encoding

Converting Decimal to Binary (2's C)

First Method: *Division*

1. Find magnitude of decimal number. (Always positive.)
2. Divide by two – remainder is least significant bit.
3. Keep dividing by two until answer is zero, writing remainders from right to left.
4. Append a zero as the MS bit;
if original number was negative, take complement and add 1.

$$X = 104_{10}$$

$$104/2 = 52 \text{ r}0 \quad \text{bit } 0$$

$$52/2 = 26 \text{ r}0 \quad \text{bit } 1$$

$$26/2 = 13 \text{ r}0 \quad \text{bit } 2$$

$$13/2 = 6 \text{ r}1 \quad \text{bit } 3$$

$$6/2 = 3 \text{ r}0 \quad \text{bit } 4$$

$$3/2 = 1 \text{ r}1 \quad \text{bit } 5$$

$$X = 01101000_2$$

$$1/2 = 0 \text{ r}1 \quad \text{bit } 6$$

Converting Decimal to Binary (2's C)

First Method: *Division*

1. Find magnitude of decimal number. (Always positive.)
2. Divide by two – remainder is least significant bit.
3. Keep dividing by two until answer is zero, writing remainders from right to left.
4. Append a zero as the MS bit;
if original number was negative, take two's complement and add 1.

$$X = -104_{10}$$

$$104/2 = 52 \text{ r}0 \quad \text{bit 0}$$

$$52/2 = 26 \text{ r}0 \quad \text{bit 1}$$

$$26/2 = 13 \text{ r}0 \quad \text{bit 2}$$

$$13/2 = 6 \text{ r}1 \quad \text{bit 3}$$

$$6/2 = 3 \text{ r}0 \quad \text{bit 4}$$

$$3/2 = 1 \text{ r}1 \quad \text{bit 5}$$

$$1/2 = 0 \text{ r}1 \quad \text{bit 6}$$

Two's complement of 1101000 = 0010111+1 = 0011000

$$X = 10011000_2$$

Converting Decimal to Binary (2's C)

Second Method: *Subtract Powers of Two*

1. Find magnitude of decimal number.
2. Subtract largest power of two less than or equal to number.
3. Put a one in the corresponding bit position.
4. Keep subtracting until result is zero.
5. Append a zero as MS bit;
if original was negative, MS bit is 1 and take complement and add 1.

<i>n</i>	2^n
0	1
1	2
2	4
3	8
4	16
5	32
6	64
7	128
8	256
9	512
10	1024

$$X = 104_{10}$$

$$104 - 64 (2^6) = 40 \quad \text{bit } 6$$

$$40 - 32 (2^5) = 8 \quad \text{bit } 5$$

$$8 - 8 (2^3) = 0 \quad \text{bit } 3$$

$$X = 01101000_2$$

Converting Decimal to Binary (2's C)

Second Method: *Subtract Powers of Two*

1. Find magnitude of decimal number.
2. Subtract largest power of two less than or equal to number.
3. Put a one in the corresponding bit position.
4. Keep subtracting until result is zero.
5. Append a zero as MS bit;
if original was negative,
MS bit is 1 and take complement and add 1.

n	2^n
0	1
1	2
2	4
3	8
4	16
5	32
6	64
7	128
8	256
9	512
10	1024

$$\begin{array}{rcll} X = -104_{10} & 104 - 64 (2^6) = 40 & \text{bit } 6 \\ & 40 - 32 (2^5) = 8 & \text{bit } 5 \\ & 8 - 8 (2^3) = 0 & \text{bit } 3 \end{array}$$

Two's complement of 1101000 = 0010111+1 = 0011000

$$X = 10011000_2$$

Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- **Integers**
 - Representation: unsigned and signed
 - **Conversion, casting**
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
 - Summary
- Representations in memory, pointers, strings

Conversions Between Signed & Unsigned

```
short int x = -12345;  
unsigned short ux = (unsigned short) x;  
printf("x = %d, ux = %u\n", x, ux);
```

Output:

x = -12345, ux = 53191

```
unsigned u = 4294967295u;    /* Umax */  
int tu = (int) u;  
printf("u = %u, tu = %d\n", u, tu);
```

Output:

u = 4294967295, tu = -1

Conversions Between Signed & Unsigned

Keep bit representations and reinterpret


Bits	Signed		Unsigned
0000	0		0
0001	1		1
0010	2		2
0011	3		3
0100	4		4
0101	5		5
0110	6	→ T2U →	6
0111	7		7
1000	-8		8
1001	-7	← U2T ←	9
1010	-6		10
1011	-5		11
1100	-4		12
1101	-3		13
1110	-2		14
1111	-1		15

Mapping Signed \leftrightarrow Unsigned

Bits	Signed		Unsigned
0000	0	\longleftrightarrow =	0
0001	1		1
0010	2		2
0011	3		3
0100	4		4
0101	5		5
0110	6		6
0111	7		7
1000	-8	\longleftrightarrow +/- 16	8
1001	-7		9
1010	-6		10
1011	-5		11
1100	-4		12
1101	-3		13
1110	-2		14
1111	-1		15

Conversion Visualized

■ 2's Complement → Unsigned



Signed vs. Unsigned in C

■ C Constants

- By **default** are considered to be **signed integers**
- **Unsigned** if has “U” as suffix

`0U, 4294967259u`

■ Casting

```
int tx, ty;
```

```
unsigned ux, uy;
```

- **Explicit casting** between signed & unsigned same as U2T and T2U

```
tx = (int) ux;          /* cast to signed */
```

```
uy = (unsigned) ty; /* cast to unsigned */
```

- **Implicit casting** also occurs via assignments and procedure calls

```
tx = ux;    /* cast to signed */
```

```
uy = ty;    /* cast to unsigned */
```

Casting Examples in C

- If there is a mix of unsigned and signed in single expression, *signed values implicitly cast to unsigned*
- Including comparison operations `<`, `>`, `==`, `<=`, `>=`
- Examples for $W = 4$: **TMin = -8**, **TMax = 7**

Constant 1	Expression	Constant 2	Type	Evaluation
0	==	0U	unsigned	1
-1	<	0	signed	1
-1	<	0U	unsigned	0
7	>	-7-1	signed	1
7U	>	-7-1	unsigned	0
-1	>	-2	signed	1
(unsigned)-1	>	-2	unsigned	1
7	>	(int) 8U	signed	1
7	<	8U	unsigned	1

Unsigned vs. Signed: Easy to Make Mistakes

```
unsigned i;  
for (i=cnt-2; i>=0; i--)      /* loop never ends */  
    a[i] += a[i+1];          /* invalid array subscript */
```

- C Standard guarantees that unsigned addition will behave like modular arithmetic
 - $0 - 1 = UMax$
- Proper way to use unsigned as loop index

```
unsigned i;  
for (i=cnt-2; i<cnt; i--)  
    a[i] += a[i+1];
```

- Even better

```
size_t i;  
for (i=cnt-2; i<cnt; i--)  
    a[i] += a[i+1];
```

- Data type **size_t** defined as unsigned value with length = word size
- Code will work even if **cnt** = *UMax*
- What if **cnt** is signed and < 0 ?

Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- **Integers**
 - Representation: unsigned and signed
 - Conversion, casting
 - **Expanding, truncating**
 - Addition, negation, multiplication, shifting
 - Summary
- Representations in memory, pointers, strings


Sign Extension

■ Task:

- Given w -bit signed integer x
- Convert it to $(w+k)$ -bit integer with same value


■ Rule:

- Make k copies of sign bit:
- $X' = \underbrace{x_{w-1}, \dots, x_{w-1}}_{k \text{ copies of MSB}}, x_{w-1}, x_{w-2}, \dots, x_0$




Sign Extension: Simple Example

Positive number



Negative number



Sign Extension Example (in C)

```
short sx = -12345;
unsigned ux = sx;           /* ux=4294954951 */
unsigned uy = (unsigned)(unsigned short) sx; /* uy=53191 */
```

	Decimal	Binary	Hex
sx	-12345	11001111 11000111	CF C7
ux	4294954951	11111111 11111111 11001111 11000111	FF FF CF C7
uy	53191	00000000 00000000 11001111 11000111	00 00 CF C7

- Converting from smaller to larger integer data type
- C automatically performs sign extension


Truncation

■ Task:

- Given $(k+w)$ -bit signed or unsigned integer X
- Convert it to w -bit integer X' with same value for “small enough” X

■ Rule:

- Drop top k bits:
- $X' = x_{w-1}, x_{w-2}, \dots, x_0$



Truncation: Simple Example

No sign change

	-16	8	4	2	1
2 =	0	0	0	1	0

	-8	4	2	1
2 =	0	0	1	0

$$2 \bmod 16 = 2$$

	-16	8	4	2	1
-6 =	1	1	0	1	0

	-8	4	2	1
-6 =	1	0	1	0

$$-6 \bmod 16 = 26 \text{U} \bmod 16 = 10 \text{U} = -6$$

Sign change

	-16	8	4	2	1
10 =	0	1	0	1	0

	-8	4	2	1
-6 =	1	0	1	0

$$10 \bmod 16 = 10 \text{U} \bmod 16 = 10 \text{U} = -6$$

	-16	8	4	2	1
-10 =	1	0	1	1	0

	-8	4	2	1
6 =	0	1	1	0

$$-10 \bmod 16 = 22 \text{U} \bmod 16 = 6 \text{U} = 6$$

Truncation Example (in C)

```
int    x = 53191;
short sx = (short) x;  /* sx = -12345 */
int    y = sx;         /* y = -12345 */
```

	Decimal	Binary	Hex
x	53191	00000000 00000000 11001111 11000111	00 00 CF C7
sx	-12345	11001111 11000111	CF C7
y	-12345	11111111 11111111 11001111 11000111	FF FF CF C7

Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- **Integers**
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - **Addition, negation, multiplication, shifting**
- Representations in memory, pointers, strings

Unsigned Addition

Operands: w bits

u

$+ v$

True Sum: $w+1$ bits

$u + v$

Discard Carry: w bits

$\text{UAdd}_w(u, v)$

■ Standard Addition Function

- Ignores carry output

■ Implements Modular Arithmetic

$$s = \text{UAdd}_w(u, v) = (u + v) \bmod 2^w$$

233	E9	1110 1001
+ 213	+ D5	+ 1101 0101
<hr/>	<hr/>	<hr/>
446	1BE	1 1011 1110
<hr/>	<hr/>	<hr/>
190	BE	1011 1110


-- discard carry

Two's Complement Addition

Operands: w bits



True Sum: $w+1$ bits



Discard Carry: w bits



■ TAdd and UAdd have Identical Bit-Level Behavior

- Signed vs. unsigned addition in C:

```
int s, t, u, v;
```

```
s = (int) ((unsigned) u + (unsigned) v);
```

```
t = u + v
```

- Will give `s == t`

```

      -23
    + -107
    -----
     -130
    -----
      126
  
```

```

      E9
    + 95
    -----
     17E
    -----
      7E
  
```

```

    1110 1001
  + 1001 0101
  -----
  1 0111 1110
  -----
    0111 1110
  
```


Negation: Complement and Increment

Operand



Take complement



Add 1



Negation: Complement and Increment

Observation: $\sim x + 1 = -x \rightarrow \sim x + x = -1$

$$\begin{array}{r} 01011100 \quad x \\ + 10100011 \quad \sim x \\ \hline 11111111 \quad -1 \end{array}$$

	Decimal	Binary	Hex
u	7	0111	7
~u	-8	1000	8
~u+1	-7	1001	9

	Decimal	Binary	Hex
u	-7	1001	9
~u	6	0110	6
~u+1	7	0111	7

Complement & Increment Examples

x=0

	Decimal	Binary	Hex
0	0	00000000 00000000	00 00
~0	-1	11111111 11111111	FF FF
~0+1	0	00000000 00000000	00 00

x=TMin

	Decimal	Binary	Hex
x	-32768	10000000 00000000	80 00
~x	32767	01111111 11111111	7F FF
~x+1	-32768	10000000 00000000	80 00


Canonical counter example

Multiplication

- **Goal: Computing Product of w -bit numbers x, y**
 - Either signed or unsigned
- **But, exact results can be bigger than w bits**
 - Up to $2w$ bits

Unsigned Multiplication

Operands: w bits



True Product: $2w$ bits



Discard w bits: w bits



$$\text{Result range: } 0 \leq x * y \leq (2^w - 1)^2 = 2^{2w} - 2^{w+1} + 1$$

■ Standard Multiplication Function

- Ignores high order w bits

■ Implements Modular Arithmetic


$$\text{UMult}_w(u, v) = (u \cdot v) \bmod 2^w$$

3	3	0011
* 5	* 5	* 0101
<hr/>	<hr/>	<hr/>
15	F	1111
<hr/>	<hr/>	<hr/>
15	F	1111

10	A	1010
* 5	* 5	* 0101
<hr/>	<hr/>	<hr/>
50	32	11 0010
<hr/>	<hr/>	<hr/>
2	2	0010

Signed Multiplication

Operands: w bits



True Product: $2w$ bits



Discard w bits: w bits



$$(-2^{w-1}) * (2^{w-1} - 1) = -2^{2w-2} + 2^{w-1} \leq x * y \leq (-2^{w-1})^2 = 2^{2w-2}$$

■ Standard Multiplication Function

- Ignores high order w bits
- Some of which are different for signed vs. unsigned multiplication
- Lower bits are the same

3	3	0011
* 5	* 5	* 0101
15	F	1111
-1	F	1111

-6	A	1010
* 5	* 5	* 0101
-30	32	11 0010
2	2	0010

Power-of-2 Multiply with Shift


■ Operation

- $u \ll k$ gives $u * 2^k$
- Both signed and unsigned

Operands: w bits

True Product: $w+k$ bits

Discard k bits: w bits



■ Examples

- $u \ll 3 == u * 8$
- $(u \ll 5) - (u \ll 3) == u * 24$
- Most machines shift and add faster than multiply
 - Compiler generates this code automatically

Compiled Multiplication Code

C Function

```
long mul_12(long x){  
    return x*12;  
}
```

Compiled Arithmetic Operations

```
leaq (%rax,%rax,2), %rax  
salq $2, %rax
```

Explanation


```
x = x + x*2  
return x << 2;
```

C compiler automatically generates shift/add code when multiplying by constant

Unsigned Power-of-2 Divide with Shift

■ Quotient of Unsigned by Power of 2

- $u \gg k$ gives $\lfloor u / 2^k \rfloor$
- Uses **logical** shift



	Division	Computed	Binary
x	15	15	1111
x >> 1	7.5	7	0 111
x >> 2	3.75	3	00 11
x >> 3	1.875	1	000 1

Compiled Unsigned Division Code

C Function

```
unsigned long udiv_8 (unsigned long x){  
    return x/8;  
}
```

Compiled Arithmetic Operations

```
shrq $3, %rax
```

Explanation


```
# Logical shift  
return x >> 3;
```

- Uses logical shift for unsigned
- For Java Users
 - Logical shift written as >>>

Signed Power-of-2 Divide with Shift

■ Quotient of Signed by Power of 2

- $x \gg k$ gives $\lfloor x / 2^k \rfloor$
- Uses **arithmetic** shift
- Rounds wrong direction if $x < 0$



	Division	Computed	Binary
x	-7	-7	1001
$x \gg 1$	-3.5	-4	1100
$x \gg 2$	-1.75	-2	1110
$x \gg 3$	-0.875	-1	1111

Compiled Signed Division Code

C Function

```
long idiv_4(long x){  
    return x/4;  
}
```

Compiled Arithmetic Operations

```
    testq %rax, %rax  
    js    L4  
L3:  
    sarq $2, %rax  
    ret  
L4:  
    addq $4, %rax  
    jmp  L3
```

Explanation


```
if x < 0  
    x += 3;  
# Arithmetic shift  
return x >> 2;
```

- Uses arithmetic shift for int
- For Java Users
 - Arith. shift written as >>

Today: Bits, Bytes, and Integers

- Representing information as bits
- Bit-level manipulations
- **Integers**
 - Representation: unsigned and signed
 - Conversion, casting
 - Expanding, truncating
 - Addition, negation, multiplication, shifting
 - Summary
- **Representations in memory, pointers, strings**

Byte-Oriented Memory Organization



- **Programs refer to data by address**
 - Conceptually, imagine it as a very large array of bytes
 - In reality, it's not, but can think of it that way
 - An address is like an index into that array
 - and, a pointer variable stores an address

- **Note: system provides private address spaces to each “process”**
 - Think of a process as a program being executed
 - So, a program can overwrite its own data, but not that of others

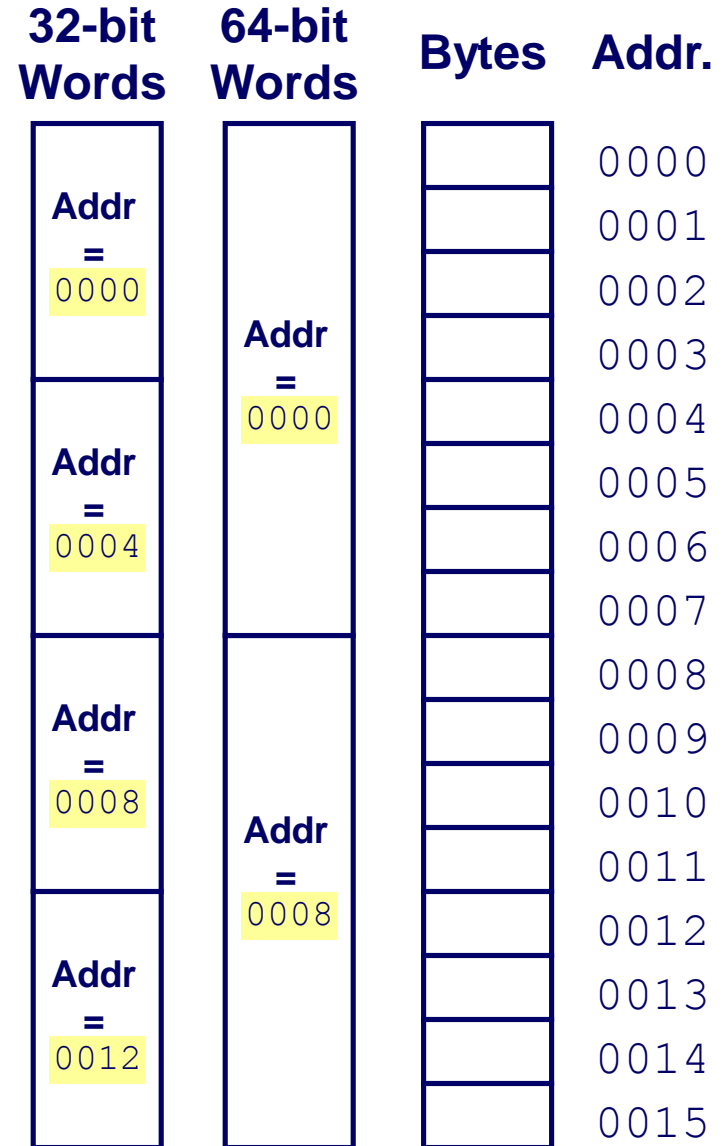
Machine Words

- **Any given computer has a “Word Size”**
 - Size of integer-valued data
 - and size of addresses
 - Until recently, most machines used 32 bits (4 bytes) as word size
 - Limits addresses to 4GB (2^{32} bytes)
 - Increasingly, machines have 64-bit word size
 - Potentially, could have 18 EB (exabytes) of addressable memory
 - That's 18.4×10^{18}
 - Machines still support multiple data formats
 - Fractions or multiples of word size
 - Always integral number of bytes

Word-Oriented Memory Organization

■ Addresses Specify Byte Locations

- Address of first byte in word
- Addresses of successive words differ by 4 (32-bit) or 8 (64-bit)



Example Data Representations

C Data Type	Typical 32-bit	Typical 64-bit	x86-64
<code>char</code>	1	1	1
<code>short</code>	2	2	2
<code>int</code>	4	4	4
<code>long</code>	4	8	8
<code>float</code>	4	4	4
<code>double</code>	8	8	8
<code>long double</code>	–	–	10/16
<code>pointer</code>	4	8	8

Byte Ordering


- So, how are the bytes within a multi-byte word ordered in memory?
- Conventions
 - Big Endian: Sun, PPC Mac, *Internet*
 - Least significant byte has **highest** address
 - Little Endian: *x86*, ARM processors running Android, iOS, and Windows
 - Least significant byte has **lowest** address

Byte Ordering Example


■ Example

- Variable `x` has 4-byte value of `0x01234567`
- Address given by `&x` is `0x100`

Big Endian



Little Endian



Representing Integers

```
int A = 12345;   long int C = 12345;
```

increasing address ↓

IA32, x86-64	Sun
39	00
30	00
00	30
00	39

IA32	x86-64	Sun
39	39	00
30	30	00
00	00	30
00	00	39
	00	
	00	
	00	
	00	

Decimal: 12345

Binary: 0011 0000 0011 1001

Hex: 3 0 3 9

Two's complement representation

```
int B = -12345;
```

IA32, x86-64	Sun
C7	FF
CF	FF
FF	CF
FF	C7

Decimal: -12345

Binary: 1100 1111 1100 0111

Hex: C F C 7

Examining Data Representations

■ Code to Print Byte Representation of Data

- Casting `pointer` to `unsigned char *` allows treatment as a byte array

```
typedef unsigned char *ucpointer;  
  
void show_bytes(ucpointer start, size_t len){  
    size_t i;  
    for (i=0; i<len; i++)  
        printf("%p\t0x%.2x\n", start+i, start[i]);  
    printf("\n");  
}
```

printf directives:

%p: Print pointer

%x: Print Hexadecimal

show_bytes Execution Example

```
int a = 12345;  
printf("int a = 12345;\n");  
show_bytes((ucpointer) &a, sizeof(int));
```

Result (Linux x86-64):

```
int a = 12345;  
0x7ffffb7f71dbc    39  
0x7ffffb7f71dbd    30  
0x7ffffb7f71dbe    00  
0x7ffffb7f71dbf    00
```

Representing Pointers

```
int B = -12345;  
int *P = &B;
```

Sun

EF
FF
FB
2C

IA32

AC
28
F5
FF

x86-64

3C
1B
FE
82
FD
7F
00
00

Different compilers & machines assign different locations to objects

Even get different results each run of the program

Representing Strings


■ Strings in C

- Represented by array of characters
- Each character encoded in ASCII format
 - Standard 7-bit encoding of character set
 - Character “0” has code 0x30
 - Digit i has code $0x30+i$
- String should be null-terminated
 - Final character = 0

■ Compatibility

- Byte ordering not an issue

```
char S[6] = "2138";
```



Reading Byte-Reversed Listings

■ Disassembly

- Text representation of binary machine code
- Generated by program that reads the machine code

■ Example Fragment (little endian machine)

Address	Instruction Code	Assembly Code
8048365:	5b	pop %ebx
8048366:	81 c3 ab 12 00 00	add \$0x12ab,%ebx
804836c:	83 bb 28 00 00 00 00	cmpl \$0x0,0x28(%ebx)

■ Deciphering Numbers

- Value: 0x12ab
- Pad to 4 bytes: 0x000012ab
- Split into bytes: 00 00 12 ab
- Reverse: ab 12 00 00

Integer C Puzzles

Initialization

```
int x = foo();  
int y = bar();  
unsigned ux = x;  
unsigned uy = y;
```

$(x < 0) \rightarrow ((x*2) < 0)$	False
$ux \geq 0$	True
$x \& 7 == 7 \rightarrow (x \ll 30) < 0$	True
$ux > -1$	False
$x > y \rightarrow -x < -y$	False
$(x*x) \geq 0$	False
$(x>0) \&\& y > 0 \rightarrow x+y > 0$	False
$x \geq 0 \rightarrow -x \leq 0$	True
$x \leq 0 \rightarrow -x \geq 0$	False
$(x -x) \gg 31 == -1$	False
$ux \gg 3 == ux/8$	True
$x \gg 3 == x/8$	False
$x \& (x-1) != 0$	False