

Machine-Level Programming I: Basics

CSE 238/2038/2138: Systems Programming

Instructor:

Fatma CORUT ERGİN

Slides adapted from Bryant & O'Hallaron's slides

Today: Machine Programming I: Basics

- History of Intel processors and architectures
- Assembly Basics: Registers, operands, move
- Arithmetic & logical operations
- C, assembly, machine code

Intel x86 Processors

- **Dominate laptop/desktop/server market**
- **Evolutionary design**
 - Backwards compatible up until 8086, introduced in 1978
 - Added more features as time goes on
 - Now 3 volumes, about 5,000 pages of documentation
- **Complex instruction set computer (CISC)**
 - Many different instructions with many different formats
 - But, only small subset encountered with Linux programs
 - Hard to match performance of Reduced Instruction Set Computers (RISC)
 - But, Intel has done just that!
 - In terms of speed. Less so for low power.

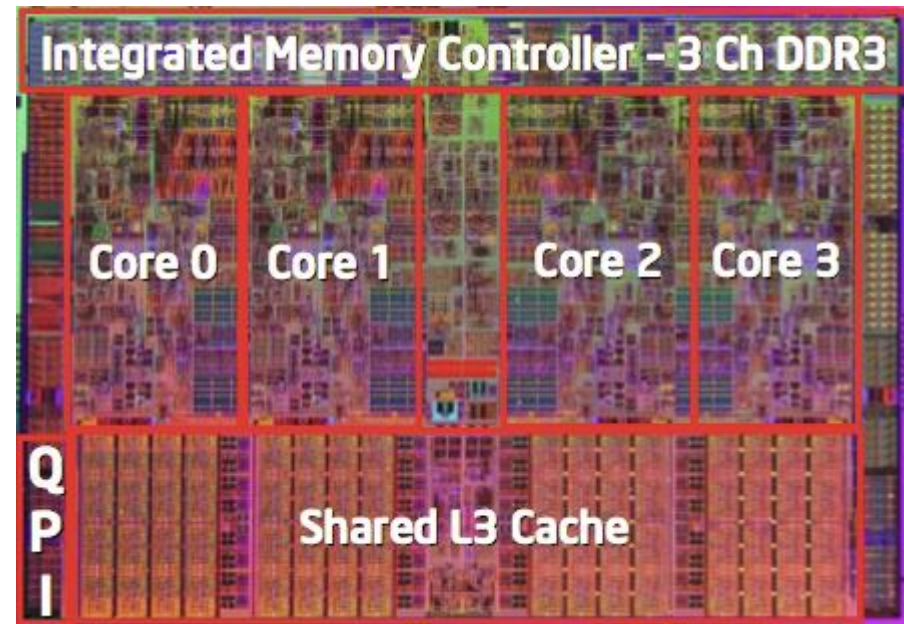
Intel x86 Evolution: Milestones

<i>Name</i>	<i>Date</i>	<i>Transistors</i>	<i>MHz</i>
■ 8086	1978	29K	5-10
<ul style="list-style-type: none">■ First 16-bit Intel processor. Basis for IBM PC & DOS■ 1MB address space			
■ 386	1985	275K	16-33
<ul style="list-style-type: none">■ First 32 bit Intel processor , referred to as IA32■ Added “flat addressing”, capable of running Unix			
■ Pentium 4E	2004	125M	2800-3800
<ul style="list-style-type: none">■ First 64-bit Intel x86 processor, referred to as x86-64			
■ Core 2	2006	291M	1060-3500
<ul style="list-style-type: none">■ First multi-core Intel processor			
■ Core i7	2008	731M	1700-3900
<ul style="list-style-type: none">■ Four cores			

Intel x86 Processors, cont.

■ Machine Evolution

■ 386	1985	0.3M
■ Pentium	1993	3.1M
■ Pentium/MMX	1997	4.5M
■ PentiumPro	1995	6.5M
■ Pentium III	1999	8.2M
■ Pentium 4	2001	42M
■ Core 2 Duo	2006	291M
■ Core i7	2008	731M
■ Core i7 Skylake	2015	1.9B



■ Added Features

- Instructions to support multimedia operations
- Instructions to enable more efficient conditional operations
- Transition from 32 bits to 64 bits
- More cores

2017 State of the Art

■ Mobile Model: Core i7

- 2.6-2.9 GHz
- 45 W

■ Desktop Model: Core i7

- Integrated graphics
- 2.8-4.0 GHz
- 35-91 W

■ Server Model: Xeon

- Integrated graphics
- Multi-socket enabled
- 2-3.7 GHz
- 25-80 W

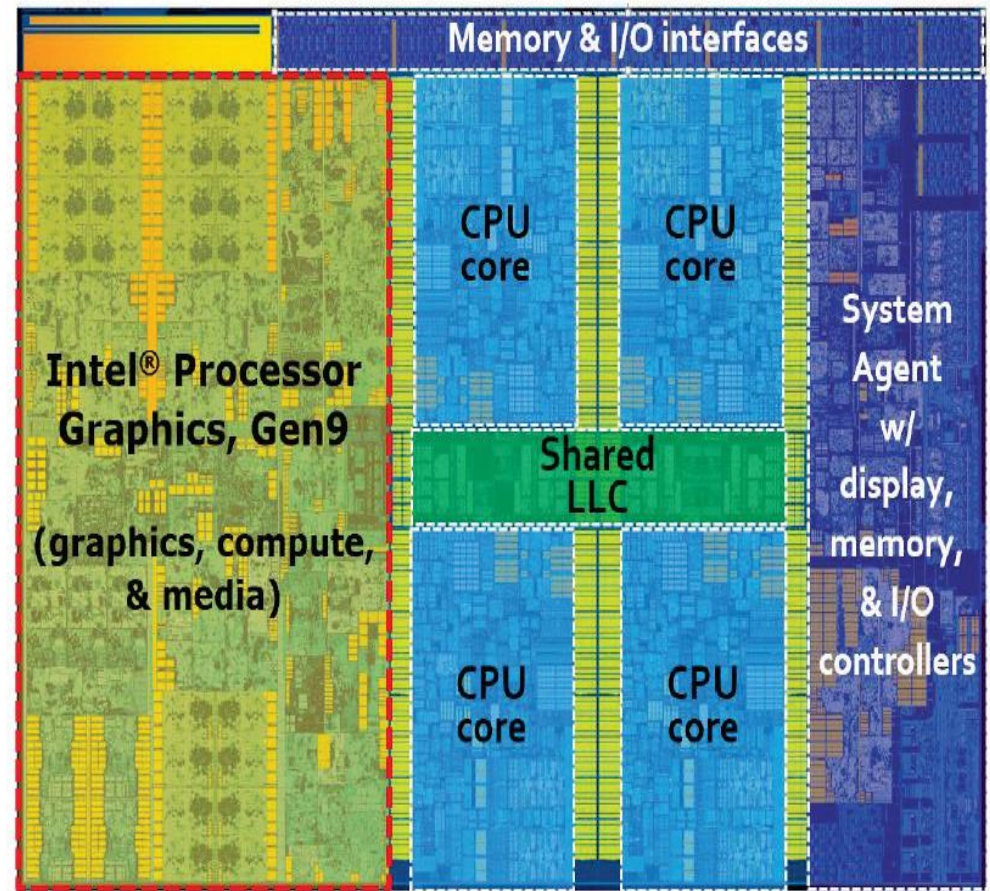


Figure 1: Architecture components layout for an Intel® Core™ i7 processor 6700K for desktop systems. This SoC contains 4 CPU cores, outlined in blue dashed boxes. Outlined in the red dashed box, is an Intel® HD Graphics 530. It is a one-slice instantiation of Intel processor graphics gen9 architecture.

x86 Clones: Advanced Micro Devices (AMD)

■ Historically

- AMD has followed just behind Intel
- A little bit slower, a lot cheaper

■ Then

- Recruited top circuit designers from Digital Equipment Corp. and other downward trending companies
- Built Opteron: tough competitor to Pentium 4
- Developed x86-64, their own extension to 64 bits

■ Recent Years

- Intel got its act together
 - Leads the world in semiconductor technology
- AMD has fallen behind
 - Relies on external semiconductor manufacturer,

Intel's 64-Bit History

- **2001: Intel Attempts Radical Shift from IA32 to IA64**
 - Totally different architecture (Itanium)
 - Executes IA32 code only as legacy
 - Performance disappointing
- **2003: AMD Steps in with Evolutionary Solution**
 - x86-64 (now called “AMD64”)
- **Intel Felt Obligated to Focus on IA64**
 - Hard to admit mistake or that AMD is better
- **2004: Intel Announces EM64T extension to IA32**
 - Extended Memory 64-bit Technology
 - Almost identical to x86-64!
- **All but low-end x86 processors support x86-64**
 - But, lots of code still runs in 32-bit mode

Our Coverage

- **x86-64**

- The standard
- `mufe> gcc hello.c`
- `mufe> gcc -m64 hello.c`


Today: Machine Programming I: Basics

- History of Intel processors and architectures
- **Assembly Basics: Registers, operands, move**
- Arithmetic & logical operations
- C, assembly, machine code

Definitions

- **Architecture:** (also ISA: instruction set architecture) The parts of a processor design that one needs to understand or write assembly/machine code.
 - Examples: instruction set specification, registers.
- **Microarchitecture:** Implementation of the architecture.
 - Examples: cache sizes and core frequency.
- **Code Forms:**
 - **Machine Code:** The byte-level programs that a processor executes
 - **Assembly Code:** A text representation of machine code
- **Example ISAs:**
 - Intel: x86, IA32, Itanium, x86-64
 - ARM: Used in almost all mobile phones
 - RISC V: New open-source ISA

Assembly/Machine Code View



Programmer-Visible State

■ PC: Program counter

- Address of next instruction
- Called “RIP” (x86-64)

■ Register file

- Heavily used program data

■ Condition codes

- Store status information about most recent arithmetic or logical operation
- Used for conditional branching

■ Memory

- Byte addressable array
- Code and user data
- Stack to support procedures

Assembly Characteristics: Data Types

- **“Integer” data of 1, 2, 4, or 8 bytes**
 - Data values
 - Addresses (untyped pointers)
- **Floating point data of 4, 8, or 10 bytes**
- **Code: Byte sequences encoding series of instructions**
- **No aggregate types such as arrays or structures**
 - Just contiguously allocated bytes in memory


x86-64 Integer Registers

%rax	%eax
%rbx	%ebx
%rcx	%ecx
%rdx	%edx
%rsi	%esi
%rdi	%edi
%rsp	%esp
%rbp	%ebp

%r8	%r8d
%r9	%r9d
%r10	%r10d
%r11	%r11d
%r12	%r12d
%r13	%r13d
%r14	%r14d
%r15	%r15d

- Can reference low-order 4 bytes (also low-order 1 & 2 bytes)
- Not part of memory (or cache)

Some History: IA32 Registers



Assembly Characteristics: Operations

- **Transfer data between memory and register**
 - Load data from memory into register
 - Store data from register into memory
- **Perform arithmetic function on register or memory data**
- **Transfer control**
 - Unconditional jumps to/from procedures
 - Conditional branches
 - Indirect branches

Moving Data

■ Moving Data

`movq Source, Dest:`

■ Operand Types

- **Immediate:** Constant integer data
 - Example: `$0x400`, `$-533`
 - Like C constant, but prefixed with ``$'`
 - Encoded with 1, 2, or 4 bytes
- **Register:** One of 16 integer registers
 - Example: `%rax`, `%r13`
 - But `%rsp` reserved for special use
 - Others have special uses for particular instructions
- **Memory:** 8 consecutive bytes of memory at address given by register
 - Simplest example: `(%rax)`
 - Various other “address modes”

`%rax`

`%rcx`

`%rdx`

`%rbx`

`%rsi`

`%rdi`

`%rsp`

`%rbp`

`%rN`

Warning: Intel docs use
`mov Dest, Source`

movq Operand Combinations

	Source	Dest	Src, Dest	C Analog
movq	Imm	Reg	movq \$0x4, %rax	temp = 0x4;
		Mem	movq \$-147, (%rax)	*p = -147;
	Reg	Reg	movq %rax, %rdx	temp2 = temp1;
		Mem	movq %rax, (%rdx)	*p = temp;
	Mem	Reg	movq (%rax), %rdx	temp = *p;

Cannot do memory-memory transfer with a single instruction

Simple Memory Addressing Modes

■ Normal (R) Mem[Reg[R]]

- Register R specifies memory address
- Pointer dereferencing in C

```
movq (%rcx), %rax
```

■ Displacement D(R) Mem[Reg[R]+D]

- Register R specifies start of memory region
- Constant displacement D specifies offset

```
movq 8(%rbp), %rdx
```

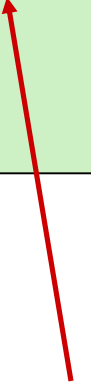
Example of Simple Addressing Modes

```
void aFunc(<type> a, <type> b) {  
    ???  
}
```

`%rdi`



`%rsi`



Example of Simple Addressing Modes

```
void swap (int *xp, int *yp){  
    int t0 = *xp;  
    int t1 = *yp;  
    *xp = t1;  
    *yp = t0;  
}
```

```
swap:  
    movl    (%rdi), %eax  
    movl    (%rsi), %ecx  
    movl    %ecx, (%rdi)  
    movl    %eax, (%rsi)  
    ret
```

Understanding Swap()

```
void swap (int *xp, int *yp){  
    int t0 = *xp;  
    int t1 = *yp;  
    *xp = t1;  
    *yp = t0;  
}
```

Register	Value
%rdi	xp
%rsi	yp
%eax	t0
%ecx	t1

Registers

%rdi	0x120	●
%rsi	0x100	●
%eax		
%ecx		

Memory

123	0x120
	0x118
	0x110
	0x108
456	0x100

swap:

```
    movl    (%rdi), %eax    # t0 = *xp  
    movl    (%rsi), %ecx    # t1 = *yp  
    movl    %ecx, (%rdi)    # *xp = t1  
    movl    %eax, (%rsi)    # *yp = t0  
    ret
```

Understanding Swap()

Registers

%rdi	0x120
%rsi	0x100
%eax	
%ecx	


Memory

Address
0x120
123
0x118
0x110
0x108
0x100
456

swap:

```
    movl    (%rdi), %eax    # t0 = *xp
    movl    (%rsi), %ecx    # t1 = *yp
    movl    %ecx, (%rdi)    # *xp = t1
    movl    %eax, (%rsi)    # *yp = t0
    ret
```


Understanding Swap()



swap:

```
movl    (%rdi), %eax    # t0 = *xp
movl    (%rsi), %ecx    # t1 = *yp
movl    %ecx, (%rdi)    # *xp = t1
movl    %eax, (%rsi)    # *yp = t0
ret
```



Understanding Swap()



swap:

```
movl    (%rdi), %eax    # t0 = *xp
movl    (%rsi), %ecx    # t1 = *yp
movl    %ecx, (%rdi)    # *xp = t1
movl    %eax, (%rsi)    # *yp = t0
ret
```

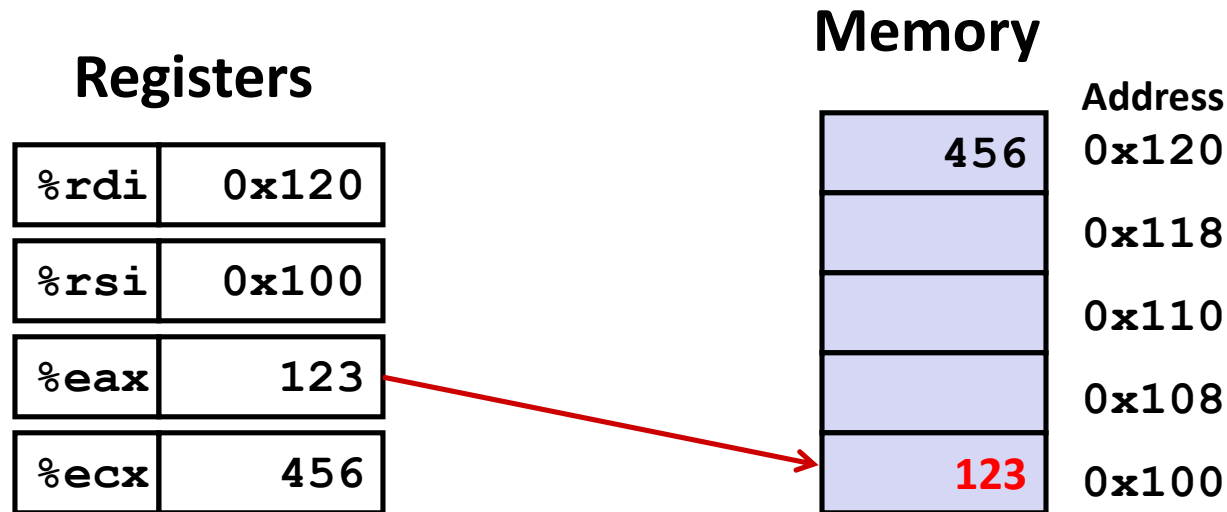
Understanding Swap()



swap:

```
movl    (%rdi), %eax    # t0 = *xp
movl    (%rsi), %ecx    # t1 = *yp
movl    %ecx, (%rdi)    # *xp = t1
movl    %eax, (%rsi)    # *yp = t0
ret
```

Understanding Swap()



swap:

```
movl    (%rdi), %eax    # t0 = *xp
movl    (%rsi), %ecx    # t1 = *yp
movl    %ecx, (%rdi)    # *xp = t1
movl    %eax, (%rsi)    # *yp = t0
ret
```

Complete Memory Addressing Modes

■ Most General Form

$D(Rb, Ri, S)$ $Mem[Reg[Rb] + S * Reg[Ri] + D]$

- D: Constant “displacement” 1, 2, or 4 bytes
- Rb: Base register: Any of 16 integer registers
- Ri: Index register: Any, except for `%rsp`
- S: Scale: 1, 2, 4, or 8 (*why these numbers?*)

■ Special Cases

(Rb, Ri) $Mem[Reg[Rb] + Reg[Ri]]$

$D(Rb, Ri)$ $Mem[Reg[Rb] + Reg[Ri] + D]$

(Rb, Ri, S) $Mem[Reg[Rb] + S * Reg[Ri]]$

Address Computation Examples

<code>%rdx</code>	<code>0xf000</code>
<code>%rcx</code>	<code>0x0100</code>

D(Rb,Ri,S)

Mem[Reg[Rb]+S*Reg[Ri]+ D]

- D: Constant “displacement” 1, 2, or 4 bytes
- Rb: Base register: Any of 16 integer registers
- Ri: Index register: Any, except for `%rsp`
- S: Scale: 1, 2, 4, or 8 (*why these numbers?*)

Expression	Address Computation	Address
<code>0x8(%rdx)</code>	<code>0xf000 + 0x8</code>	<code>0xf008</code>
<code>(%rdx,%rcx)</code>	<code>0xf000 + 0x100</code>	<code>0xf100</code>
<code>(%rdx,%rcx,4)</code>	<code>0xf000 + 4*0x100</code>	<code>0xf400</code>
<code>0x80(,%rdx,2)</code>	<code>2*0xf000 + 0x80</code>	<code>0x1e080</code>

Today: Machine Programming I: Basics

- History of Intel processors and architectures
- Assembly Basics: Registers, operands, move
- **Arithmetic & logical operations**
- C, assembly, machine code

Address Computation Instruction

■ `leaq Src, Dst`

- *Src* is address mode expression
- Set *Dst* to address denoted by expression

■ Uses

- Computing addresses without a memory reference
 - E.g., translation of `p = &x[i];`
- Computing arithmetic expressions of the form $x + k * y$
 - $k = 1, 2, 4, \text{ or } 8$

■ Example

```
long m12(long x){  
    return x*12;  
}
```

Converted to ASM by compiler:

```
leaq (%rdi,%rdi,2), %rax # t <- x+x*2  
salq $2, %rax           # return t<<2
```

Some Arithmetic Operations

■ Two Operand Instructions:

<i>Format</i>	<i>Computation</i>	
<code>addq Src, Dest</code>	$\text{Dest} = \text{Dest} + \text{Src}$	
<code>subq Src, Dest</code>	$\text{Dest} = \text{Dest} - \text{Src}$	
<code>imulq Src, Dest</code>	$\text{Dest} = \text{Dest} * \text{Src}$	
<code>salq Src, Dest</code>	$\text{Dest} = \text{Dest} \ll \text{Src}$	<i>Also called <code>shlq</code></i>
<code>sarq Src, Dest</code>	$\text{Dest} = \text{Dest} \gg \text{Src}$	<i>Arithmetic right shift</i>
<code>shrq Src, Dest</code>	$\text{Dest} = \text{Dest} \gg \text{Src}$	<i>Logical right shift</i>
<code>xorq Src, Dest</code>	$\text{Dest} = \text{Dest} \wedge \text{Src}$	
<code>andq Src, Dest</code>	$\text{Dest} = \text{Dest} \& \text{Src}$	
<code>orq Src, Dest</code>	$\text{Dest} = \text{Dest} \text{Src}$	

■ Watch out for argument order!

(Warning: Intel docs use “op Dest,Src”)

■ No distinction between signed and unsigned int (why?)

Some Arithmetic Operations

■ One Operand Instructions

<code>incq</code>	<code>Dest</code>	$\text{Dest} = \text{Dest} + 1$
<code>decq</code>	<code>Dest</code>	$\text{Dest} = \text{Dest} - 1$
<code>negq</code>	<code>Dest</code>	$\text{Dest} = -\text{Dest}$
<code>notq</code>	<code>Dest</code>	$\text{Dest} = \sim\text{Dest}$

■ See book for more instructions

Arithmetic Expression Example

```
long arith
(long x, long y, long z){
    long t1 = x+y;
    long t2 = z+t1;
    long t3 = x+4;
    long t4 = y * 48;
    long t5 = t3 + t4;
    long rval = t2 * t5;
    return rval;
}
```

```
arith:
    leaq    (%rdi,%rsi), %rax
    addq    %rdx, %rax
    leaq    (%rsi,%rsi,2), %rdx
    salq    $4, %rdx
    leaq    4(%rdi,%rdx), %rcx
    imulq    %rcx, %rax
    ret
```

Interesting Instructions

- **leaq**: address computation
- **salq**: shift
- **imulq**: multiplication
 - But, only used once

Understanding Arithmetic Expression

Example

```
long arith
(long x, long y, long z){
    long t1 = x+y;
    long t2 = z+t1;
    long t3 = x+4;
    long t4 = y * 48;
    long t5 = t3 + t4;
    long rval = t2 * t5;
    return rval;
}
```

```
arith:
    leaq    (%rdi,%rsi), %rax    # t1
    addq    %rdx, %rax          # t2
    leaq    (%rsi,%rsi,2), %rdx
    salq    $4, %rdx            # t4
    leaq    4(%rdi,%rdx), %rcx   # t5
    imulq    %rcx, %rax          # rval
    ret
```

Register	Use(s)
%rdi	Argument x
%rsi	Argument y
%rdx	Argument z , t4
%rax	t1 , t2 , rval
%rcx	t5

Compiler optimization:

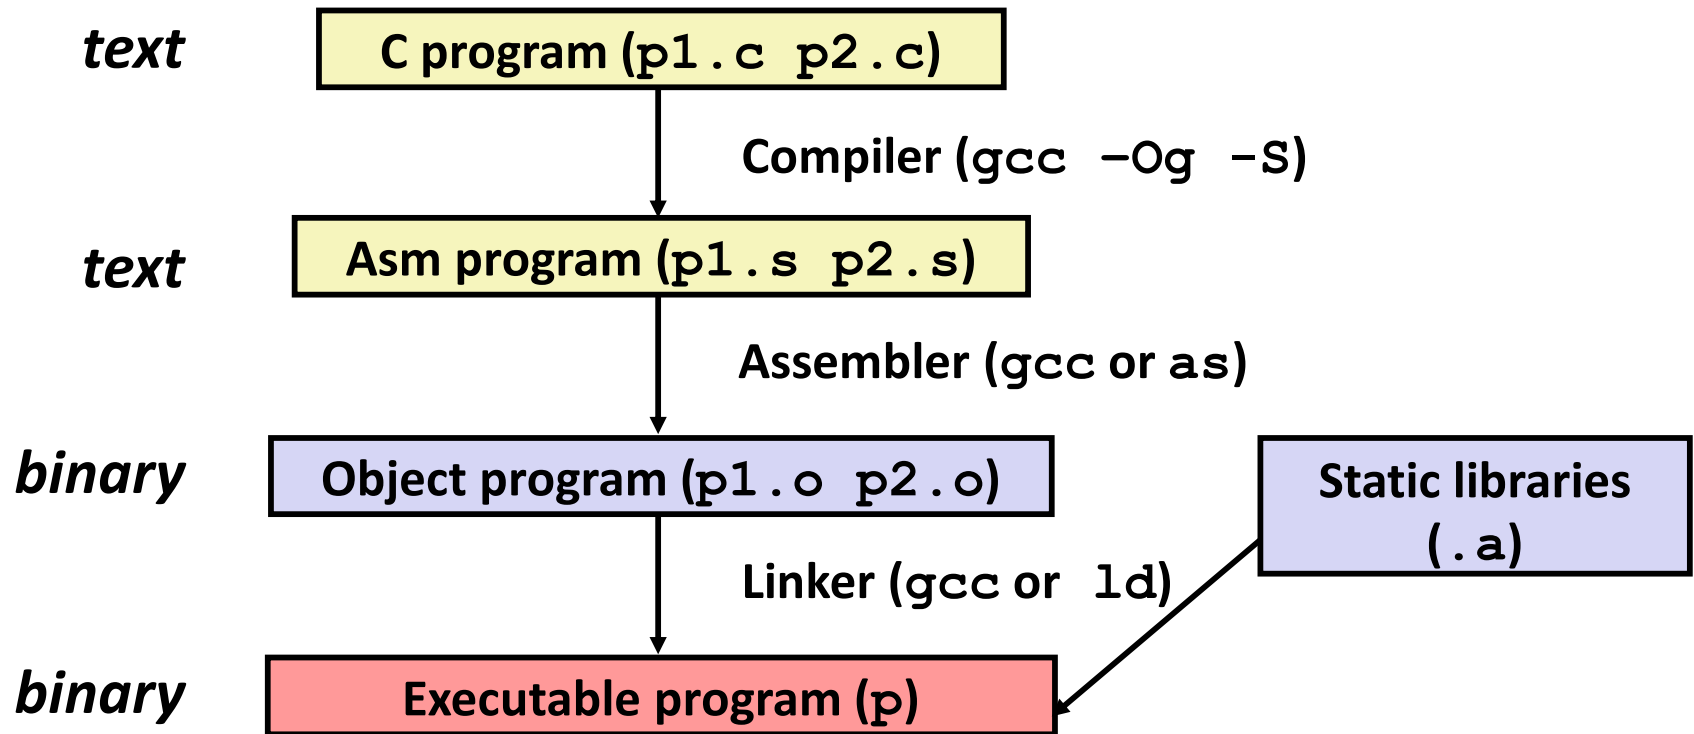
- Reuse of registers
- Substitution (copy propagation)
- Strength reduction

Today: Machine Programming I: Basics

- History of Intel processors and architectures
- Assembly Basics: Registers, operands, move
- Arithmetic & logical operations
- **C, assembly, machine code**

Turning C into Object Code

- Code in files `p1.c` `p2.c`
- Compile with command: `gcc -Og p1.c p2.c -o p`
 - Use basic optimizations (`-Og`) [New to recent versions of GCC]
 - Put resulting binary in file `p`



Compiling Into Assembly

C Code (sum.c)

```
int plus(int x, int y){  
    return x+y;  
}  
void sumstore(int x, int y, int *dest){  
    int t = plus(x, y);  
    *dest = t;  
}
```

Generated x86-64 Assembly

```
sumstore:  
    pushq    %rbp  
    movq     %rsp, %rbp  
    addl     %esi, %edi  
    movl     %edi, (%rdx)  
    popq     %rbp  
    ret
```

Obtain (on Mac OS) with command

```
gcc -O -S sum.c
```

Produces file `sum.s`

Warning: Will get very different results on different machines (Linux, Mac OS-X, ...) due to different versions of gcc and different compiler settings.

What it Really Looks Like

```
.globl _sumstore
.align 4, 0x90
_sumstore:
.cfi_startproc
## BB#0:
    pushq    %rbp

Ltmp3:
    .cfi_def_cfa_offset 16

Ltmp4:
    .cfi_offset %rbp, -16
    movq     %rsp, %rbp

Ltmp5:
    .cfi_def_cfa_register %rbp
    addl     %esi, %edi
    movl     %edi, (%rdx)
    popq     %rbp
    retq
.cfi_endproc
```

Things preceding with a "." are generally directives

```
sumstore:
    pushq    %rbp
    movq     %rsp, %rbp
    addl     %esi, %edi
    movl     %edi, (%rdx)
    popq     %rbp
    ret
```

Object Code

Code for `sumstore`

`0x100000f20:`

`0x55`

`0x48`

`0x89`

`0xe5`

`0x01`

`0xf7`

`0x89`

`0x3a`

`0x5d`

`0xc3`

- Total of 10 bytes
- Each instruction 1, 2, or 3 bytes
- Starts at address `0x100000f20`

■ Assembler

- Translates `.s` into `.o`
- Binary encoding of each instruction
- Nearly-complete image of executable code
- Missing linkages between code in different files

■ Linker

- Resolves references between files
- Combines with static run-time libraries
 - E.g., code for `malloc`, `printf`
- Some libraries are *dynamically linked*
 - Linking occurs when program begins execution

Machine Instruction Example

```
*dest = t;
```

```
movl    %edi, (%rdx)
```

```
0x100000f26: 89 3a
```

■ C Code

- Store value **t** where designated by **dest**

■ Assembly

- Move 4-byte value to memory
- Operands:
 - t**: Register **%edi**
 - dest**: Register **%rdx**
 - *dest**: Memory **M[%rdx]**

■ Object Code

- 2-byte instruction
- Stored at address **0x100000f26**

Disassembling Object Code

■ Disassembler

`objdump -d sum`

- Useful tool for examining object code
- Analyzes bit pattern of series of instructions
- Produces approximate rendition of assembly code
- Can be run on either a `.out` (complete executable) or `.o` file

Disassembled

`<sumstore>:`

<code>100000f20:</code>	<code>55</code>	<code>pushq</code>	<code>%rbp</code>
<code>100000f21:</code>	<code>48 89 e5</code>	<code>movq</code>	<code>%rsp, %rbp</code>
<code>100000f24:</code>	<code>01 f7</code>	<code>addl</code>	<code>%esi, %edi</code>
<code>100000f26:</code>	<code>89 3a</code>	<code>movl</code>	<code>%edi, (%rdx)</code>
<code>100000f28:</code>	<code>5d</code>	<code>popq</code>	<code>%rbp</code>
<code>100000f29:</code>	<code>c3</code>	<code>retq</code>	

Alternate Disassembly

Object

0x100000f20:

0x55

0x48

0x89

0xe5

0x01

0xf7

0x89

0x3a

0x5d

0xc3

Disassembled

Dump of assembler code for function sumstore:

```
0x00000000100000f20 <+0>: pushq    %rbp
0x00000000100000f21 <+1>: movq    %rsp,%rbp
0x00000000100000f24 <+4>: addl    %esi, %edi
0x00000000100000f26 <+6>: movl    %edi, (%rdx)
0x00000000100000f28 <+8>: popq    %rbp
0x00000000100000f29 <+9>: retq
```

■ Within gdb Debugger

- `gdb sum`
- `disassemble sumstore`
 - Disassemble procedure
- `x/10xb sumstore`
 - Examine the 10 bytes starting at `sumstore`

What Can be Disassembled?

```
% objdump -d WINWORD.EXE
```

```
WINWORD.EXE:      file format pei-i386
```

```
No symbols in "WINWORD.EXE".
```

```
Disassembly of section .text:
```

```
30001000 <.text>:
```

```
30001000:
```

```
30001001:
```

```
30001003:
```

```
30001005:
```

```
3000100a:
```

**Reverse engineering forbidden by
Microsoft End User License Agreement**

- Anything that can be interpreted as executable code
- Disassembler examines bytes and reconstructs assembly source

Machine Programming I: Summary

■ History of Intel processors and architectures

- Evolutionary design leads to many quirks and artifacts

■ C, assembly, machine code

- New forms of visible state: program counter, registers, ...
- Compiler must transform statements, expressions, procedures into low-level instruction sequences

■ Assembly Basics: Registers, operands, move

- The x86-64 move instructions cover wide range of data movement forms

■ Arithmetic

- C compiler will figure out different instruction combinations to carry out computation