"On my honor, I have neither given nor received any unauthorized and/or inappropriate assistance for all sessions of this exam. The work done on this exam is totally my own. I understand that by the school code, violation of these principles will lead to a zero grade and is subject to harsh discipline issue."

Huseyin Kerem Mican

15011 96 29

Q1) int -> 4 bytes   pointer   imul -> çarpım

M -> 18

N -> 12

int i , int j , A[M][N] , int B[N][M]
(addr over A[M][N])

return A[i][j] + B[i][j]

%rsp , %rbp

%rdx , %rdi

leaq (%rax, %rax, 4) (4 = scale value)

8+4 — 12

Q2)

union -> 8 char * -> 8 bytes

1) char * one(struct STR *ptr) { return ptr->un.d; } → pointer will gave address

c) movq 8(%rdi), %rax
movsbl (%rax), %eax

1 → c

2) " " return ptr -> next -> un.f[3]

d) movq 16(%rdi), %rax
movsbl (%rax), %eax

2 → d

3) return ptr -> b → short

a) movswl 2(%rdi), %eax)

3 → a

4) return *(ptr -> un.d)

b) movq 8(%rdi), %eax

4 → b

5. (12 pts) The three functions below perform the same operation with varying degrees of spatial locality. Rank-order the functions with respect to the spatial locality enjoyed by each. Also, write an explanation for your ranking.

```
#define n 100
typedef struct {
    int x[3];
    int y[3];
} pt;

pt p[n];
```

```
void f1(pt *p, int n){
    int i, j;
    for (j=0; j<3; j++){
        for (i=0; i<n; i++)
            p[i].x[j] = 0;
        for (i=0; i<n; i++)
            p[i].y[j] = 0;
    }
}
```

```
void f2(pt *p, int n){
    int i, j;
    for (i=0; i<n; i++){
        for (j=0; j<3; j++)
            p[i].x[j] = 0;
        for (j=0; j<3; j++)
            p[i].y[j] = 0;
    }
}
```

```
void f3(pt *p, int n){
    int i, j;
    for (i=0; i<n; i++){
        for (j=0; j<3; j++)
            p[i].x[j] = 0;
            p[i].y[j] = 0;
    }
}
```

a) best spatial locality f2

b) mid spatial locality f3

c) worst spatial locality f1

pt+0 x[3] pt+12 y[3] pt+24 mult. of 4

f2 (B) f3 (B) f1

p[0] { p[0].x[0], p[0].x[1], p[0].x[2], p[0].y[0], p[0].y[1], p[0].y[2] }
p[1] { p[1].x[0], p[1].x[1], p[1].x[2], p[1].y[0] }

f1: p[0].x[0], p[1].x[0], p[2].x[0] ... p[99].x[0], p[0].y[0], p[1].y[0], p[2].y[0] ... p[99].y[0]

- Physical addresses are 8 bits long
- The block size is 4 bytes (B = 4)
- The cache has 4 sets (S = 4)
- The cache is direct mapped (E = 1)

a) What is the total capacity of the cache? (in number of data bytes)

b) How long is a tag? (in number of bits) 4 bits

c) Assuming that the cache starts clean (all lines invalid), please fill in the following tables, describing what happens with each operation. Addresses are given in both hex and binary for your convenience.

SI	V	tag	Blocks
00	1	0010	20, 21, 22, 23
01	1	0000	04, 05, 06, 07
10	1	0001	18, 19, 1A, 1B
11	0		

Operation	Set index?	Hit or Miss?	Eviction?
load 0x00 (0000 0000)₂	0	Miss	no
load 0x04 (0000 0100)₂	1	Miss	no
load 0x08 (0000 1000)₂	2	Miss	no
store 0x12 (0001 0010)₂	0	Miss	yes
load 0x16 (0001 0110)₂	1	Miss	yes
store 0x06 (0000 0110)₂	1	Miss	yes
load 0x18 (0001 1000)₂	2	Miss	yes
load 0x20 (0010 0000)₂	0	Miss	yes
store 0x1A (0001 1010)₂	2	Hit	no

tag | SI | BO

0x 00 01 02 03 04 05 06 07

00000100, 00000101, 00000110, 00000111

B. 00000000, 00000001, 00000010, 00000011

---

v. Give the contents of header and footer for each block (in binary). Don't forget to indicate the change in header and footer when there is a need.

a) p1 = malloc(3)
b) p2 = malloc(5)
c) free(p1)
d) p3 = malloc(8)
e) free(p2)
f) p4 = malloc(4)

Show all your work and briefly explain your answer!

5+2+2 = 9 Payload + H + F; 12 bytes

3 bytes + (H) + (F) (3)+(2)+(2) (7) ← 4-byte alignment; Payload + H + F = 8 bytes

00.. 001111   0000000000000001001   0.. 001001

9+2+2 = 1[illegible]

The heap will be like the following, and you will draw the heap 5 times (i.e., for each request). You can insert more rows as you need.

Address	Content	Header/Footer?	is pointed?
0xb000	0000000000001001 Header		
0xb002	Free		
0xb004	Free		
0xb006	0000000000001001 Footer		
0xb008	0000000000001111 Header		
0xb00a	Allocated		p2 points at this m.l.
0xb00c	Allocated		
0xb00e	Allocated		
0xb010	Allocated		
	00.. 01111 Footer		

Here is an example of what you will draw to show your answer. Note that this example is for a totally different allocation.

Address	Content	Header/Footer?	is pointed?
0x400AFFC	0x00000013	Header	
0x4008000	Allocated		p1 points at this memory location
0x400b004	Allocated		
0x400b008	0x00000013	Footer	

Array Example

Practice Problem 3.38 (solution page 377)

Consider the following source code, where *M* and *N* are constants declared with #define:

```
long P[M][N];
long Q[N][M];

long sum_element(long i, long j) {
    return P[i][j] + Q[j][i];
}
```

```
long sum_element(long i, long j)
i in %rdi, j in %rsi
sum_element:
  leaq    0(,%rdi,8), %rdx      Compute 8i → rdx
  subq    %rdi, %rdx            Compute 7i → rdx
  addq    %rsi, %rdx            Compute 7i + j → rdx
  leaq    (%rsi,%rsi,4), %rax   Compute 5j → rax
  addq    %rax, %rdi            Compute i + 5j → rdi
  movq    Q(,%rdi,8), %rax      Retrieve M[x_Q+8*(5j + i)] → rax
  addq    P(,%rdx,8), %rax      Add M[x_P + 8*(7i+j)] → rax
  ret
```

Use your reverse engineering skills to determine the values of *M* and *N* based on this assembly code.

8 is scale value

Offset of matrix P 7i + j →
P has 7 columns

Offset of matrix Q 5j + i →
Q has 5 columns

So, M = 5 and N = 7

General Cache Organization (S, E, B)

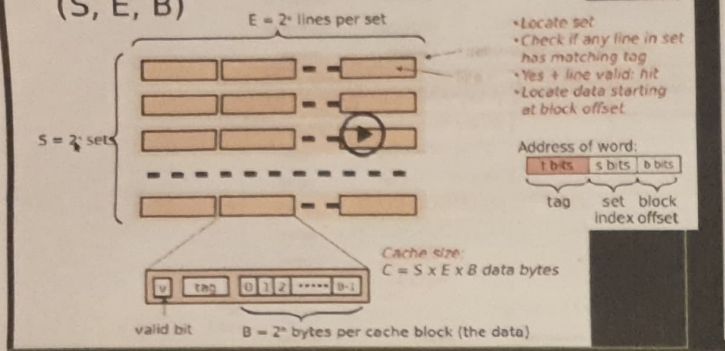

Parameter	Description
Fundamental parameters	
$S = 2^s$	Number of sets
E	Number of lines per set
$B = 2^b$	Block size (bytes)
$m = \log_2(M)$	Number of physical (main memory) address bits
Derived quantities	
$M = 2^m$	Maximum number of unique memory addresses
$s = \log_2(S)$	Number of *set index bits*
$b = \log_2(B)$	Number of *block offset bits*
$t = m - (s + b)$	Number of *tag bits*
$C = B \times E \times S$	Cache size (bytes), not including overhead such as the valid and tag bits

What about writes?

- **Multiple copies of data exist:**
  - L1, L2, L3, Main Memory, Disk
- **What to do on a write-hit?**
  - Write-through (write immediately to memory)
  - Write-back (defer write to memory until replacement of line)
    - Need a dirty bit (line different from memory or not)
- **What to do on a write-miss?**
  - Write-allocate (load into cache, update line in cache)
    - Good if more writes to the location follow
  - No-write-allocate (writes straight to memory, does not load into cache)
- **Typical**
  - Write-through + No-write-allocate
  - **Write-back + Write-allocate**

Satisfying Alignment with Structures

- **Within structure:**
  - Must satisfy each element's alignment requirement
- **Overall structure placement**
  - Each structure has alignment requirement K
    - K = Largest alignment of any element
  - Initial address & structure length must be multiples of K
- **Example:**
  - K = 8, due to double element

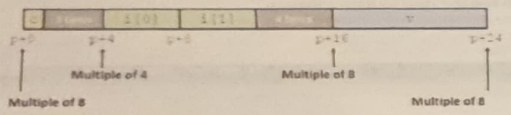

Union Allocation

- **Allocate according to largest element**
- **Can only use one field at a time**

```
union U1 {
  char c;
  int i[2];
  double v;
} *up;
```

```
struct S1 {
  char c;
  int i[2];
  double v;
} *sp;
```

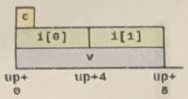

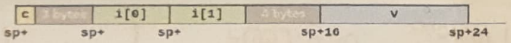

FP Basics

- **Arguments passed in %xmm0, %xmm1, ...**
- **Result returned in %xmm0**
- **All XMM registers caller-saved**

```
float fadd(float x, float y)
{
    return x + y;
}
```

```
# x in %xmm0, y in %xmm1
addss   %xmm1, %xmm0
ret
```

```
double dadd(double x, double y)
{
    return x + y;
}
```

```
# x in %xmm0, y in %xmm1
addsd   %xmm1, %xmm0
ret
```

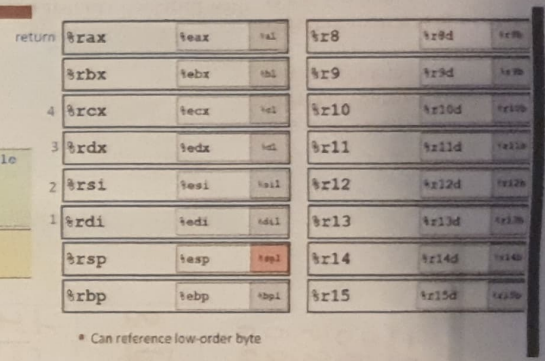

* Can reference low-order byte

Locality

- **Principle of Locality:** Programs tend to use data and instructions with addresses near or equal to those they have used recently
- **Temporal locality:**
  - Recently referenced items are likely to be referenced again in the near future
- **Spatial locality:**
  - Items with nearby addresses tend to be referenced close together in time

General Caching Concepts: Types of Cache Misses

- **Cold (compulsory) miss**
  - Cold misses occur because the cache is empty.
- **Conflict miss**
  - Most caches limit blocks at level k+1 to a small subset (sometimes a singleton) of the block positions at level k.
    - E.g. Block i at level k+1 must be placed in block (i mod 4) at level k.
  - Conflict misses occur when the level k cache is large enough, but multiple data objects all map to the same level k block.
    - E.g. Referencing blocks 0, 8, 0, 8, 0, 8, ... would miss every time.
- **Capacity miss**
  - Occurs when the set of active cache blocks (working set) is larger than the cache.

Example: Direct Mapped Cache (E = 1)

Direct mapped: One line per set
Assume: cache block size 8 bytes

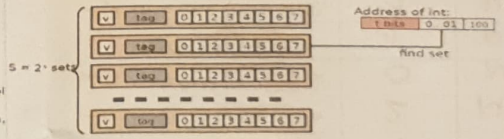

If tag doesn't match: old line is evicted and replaced

E-way Set Associative Cache (Here: E = 2)

E = 2: Two lines per set
Assume: cache block size 8 bytes

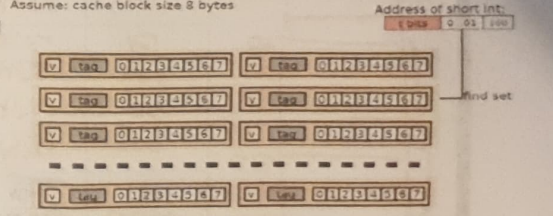

Keeping Track of Free Blocks

- **Method 1: *Implicit list* using length—links all blocks**

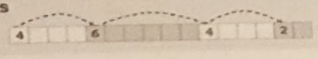

Need to tag each block as allocated/free

- **Method 2: *Explicit list* among the free blocks using pointers**

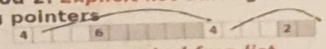

Need space for pointers

- **Method 3: *Segregated free list***
  - Different free lists for different size classes
- **Method 4: *Blocks sorted by size***
  - Can use a balanced tree (e.g. Red-Black tree) with pointers within each free block, and the length used as a key

Method 1: Implicit List

- **For each block we need both size and allocation status**
  - Could store this information in two words: wasteful!
- **Standard trick**
  - If blocks are aligned, some low-order address bits are always 0
  - Instead of storing an always-0 bit, use it as a allocated/free flag
  - When reading size word, must mask out this bit

Format of allocated and free blocks

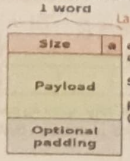

Implicit List: Finding a Free Block

- ***First fit:***
  - Search list from beginning, choose *first* free block that fits
  - Can take linear time in total number of blocks (allocated and free)
  - In practice it can cause "splinters" at beginning of list
- ***Next fit:***
  - Like first fit, but search list starting where previous search finished
  - Should often be faster than first fit: avoids re-scanning unhelpful blocks
  - Some research suggests that fragmentation is worse
- ***Best fit:***
  - Search the list, choose the *best* free block: fits, with fewest bytes left over
  - Keeps fragments small—usually improves memory utilization
  - Will typically run slower than first fit

Nested Array Element Access

- Array Elements
- `A[I][J]` is element of type T, which requires K bytes
- Address `A + I * (C * K) + J * K`
  `= A + (I * C + J) * K`

`int A[R][C];`