




The image cannot be displayed. Your computer may not have enough memory to open the image, or the image may have been corrupted. Restart your computer, and then open the file again. If the red x still appears, you may have to delete the image and then insert it again.

Threads

Chapter 4

Definition of Threads

- Basic unit of CPU utilization.
- Execution context that is independently scheduled but shares a single address space with other threads.
- Traditional Process = Single threaded process (single thread of execution per process – The concept of thread not recognized)
- Multithreaded Process = Multiple threads of execution.



What is associated with a process?

- A virtual address space which holds the process image
 - It includes:
 - ▶ PCB,
 - ▶ user-address space (data & codes segments),
 - ▶ user and kernel stack
 - ▶ OS resources (open files & signals)
- Protected access to processor, other processes (IPC), files & I/O resources.

Thread ?

- An execution state (running, ready, etc.)
- Saved thread context when not running
- **Separate for each thread**
 - Thread ID, PC, register set, scheduling properties
(TCB = thread control block)
 - User and kernel stack (execution stack; some per-thread static storage for local variables)
- **Shared Among Threads**
 - Access to the memory and resources of its process
 - User address space (data & code segment)
 - OS resources (open files & signals)

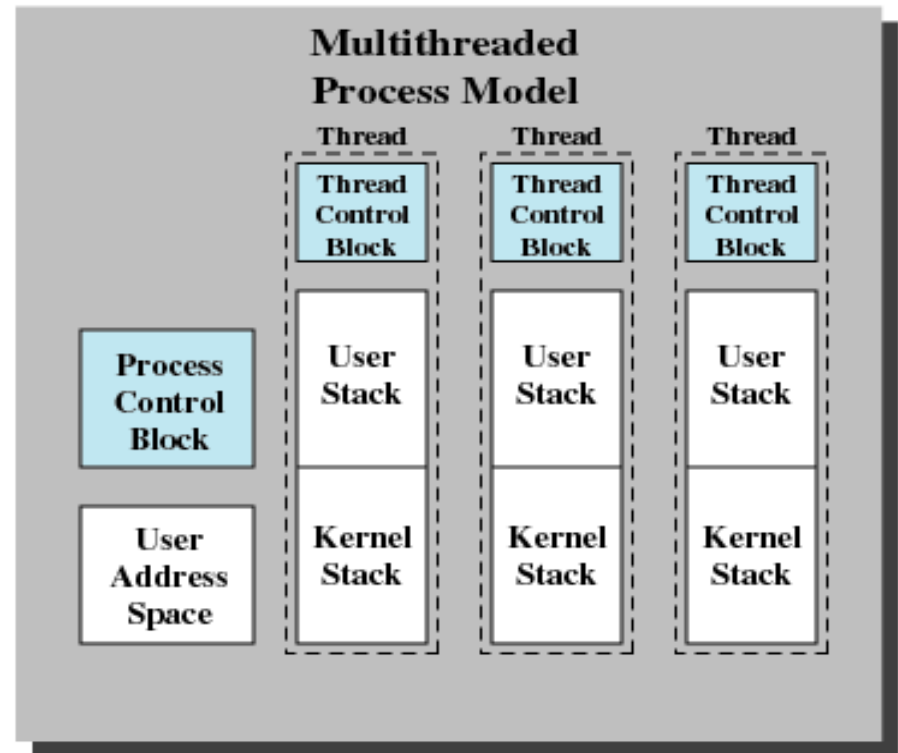
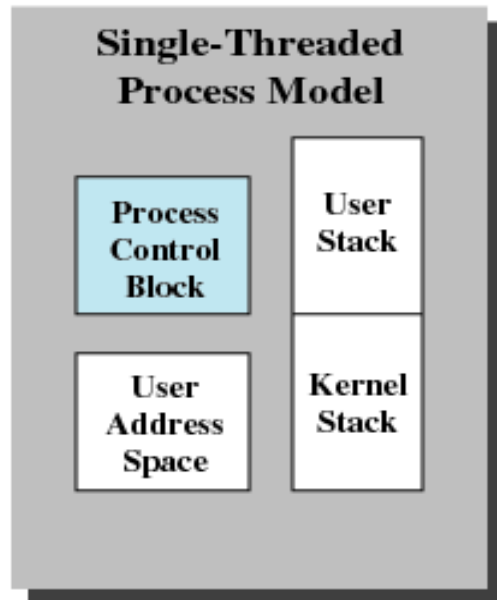


Figure 4.2 Single Threaded and Multithreaded Process Models

Benefits of Threads

1. Economy

- Takes less time to create a new thread than a process
- Less time to terminate a thread than a process
- Less time to switch between two threads within the same process
- Since threads within the same process share memory and files, they can communicate with each other without invoking the kernel

Benefits of Threads (2)

2. Responsiveness

- allow a program to continue running even if part of it blocked

3. Resource Sharing

- Share memory and resources of the process

Examples (1)

- Foreground to background work
 - **web browser** = one thread display image / text; another thread retrieve data from network
 - **spreadsheet program** = one thread display menu & read user input; another execute user command and update the spreadsheet
 - **web server** = when it receives a request:
 - ▶ create a separate process to service the request (extra overhead)
 - ▶ Solution: multiple threads to serve the same purpose

Examples (2)

■ Asynchronous processing

- **Word processor**: write its RAM buffer to disk once every minute.

■ Speed of execution

- Compute one batch of data while reading the next batch from a device.

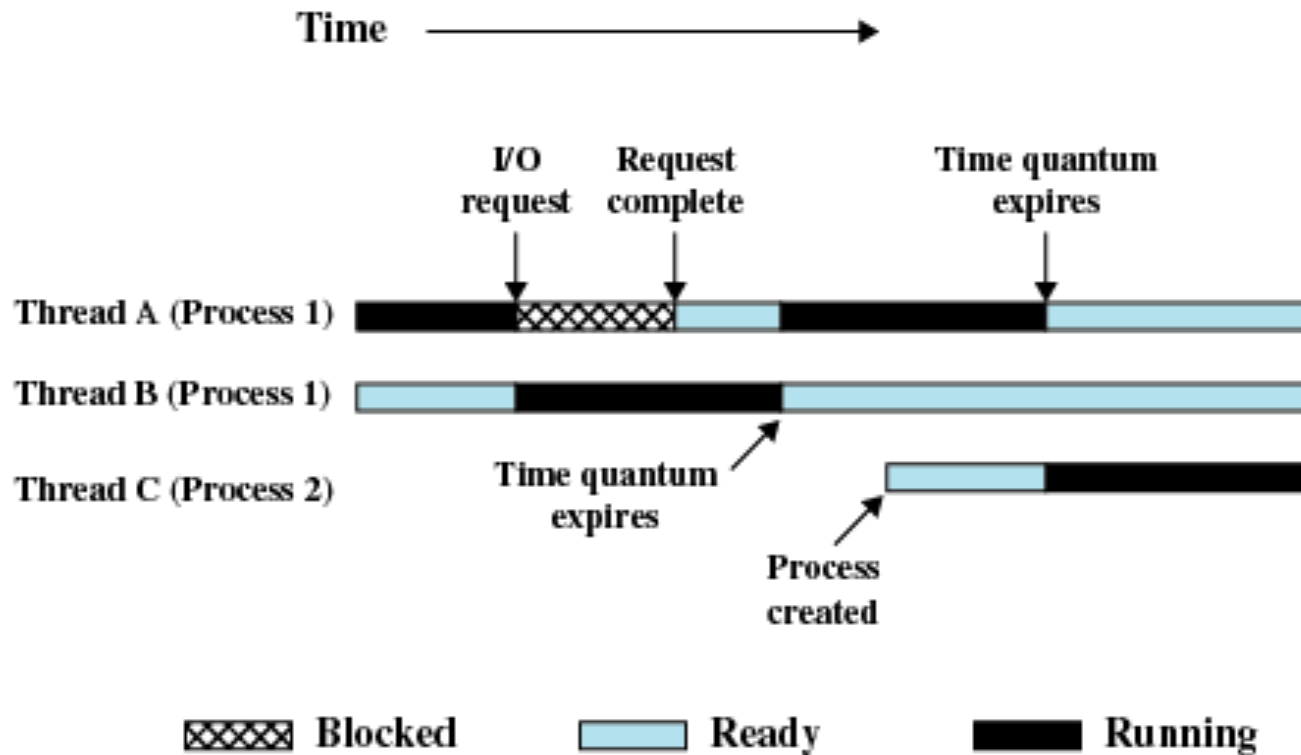
Suspension and Termination of Threads

- Suspending a process involves suspending all threads of the process since all threads share the same address space
- Termination of a process, terminates all threads within the process

Thread States


- States (running, ready, blocked)
- Basic thread operations associated with a change in thread state
 - Spawn (spawn another thread)
 - Block
 - Unblock
 - Finish (deallocate register context and stacks)

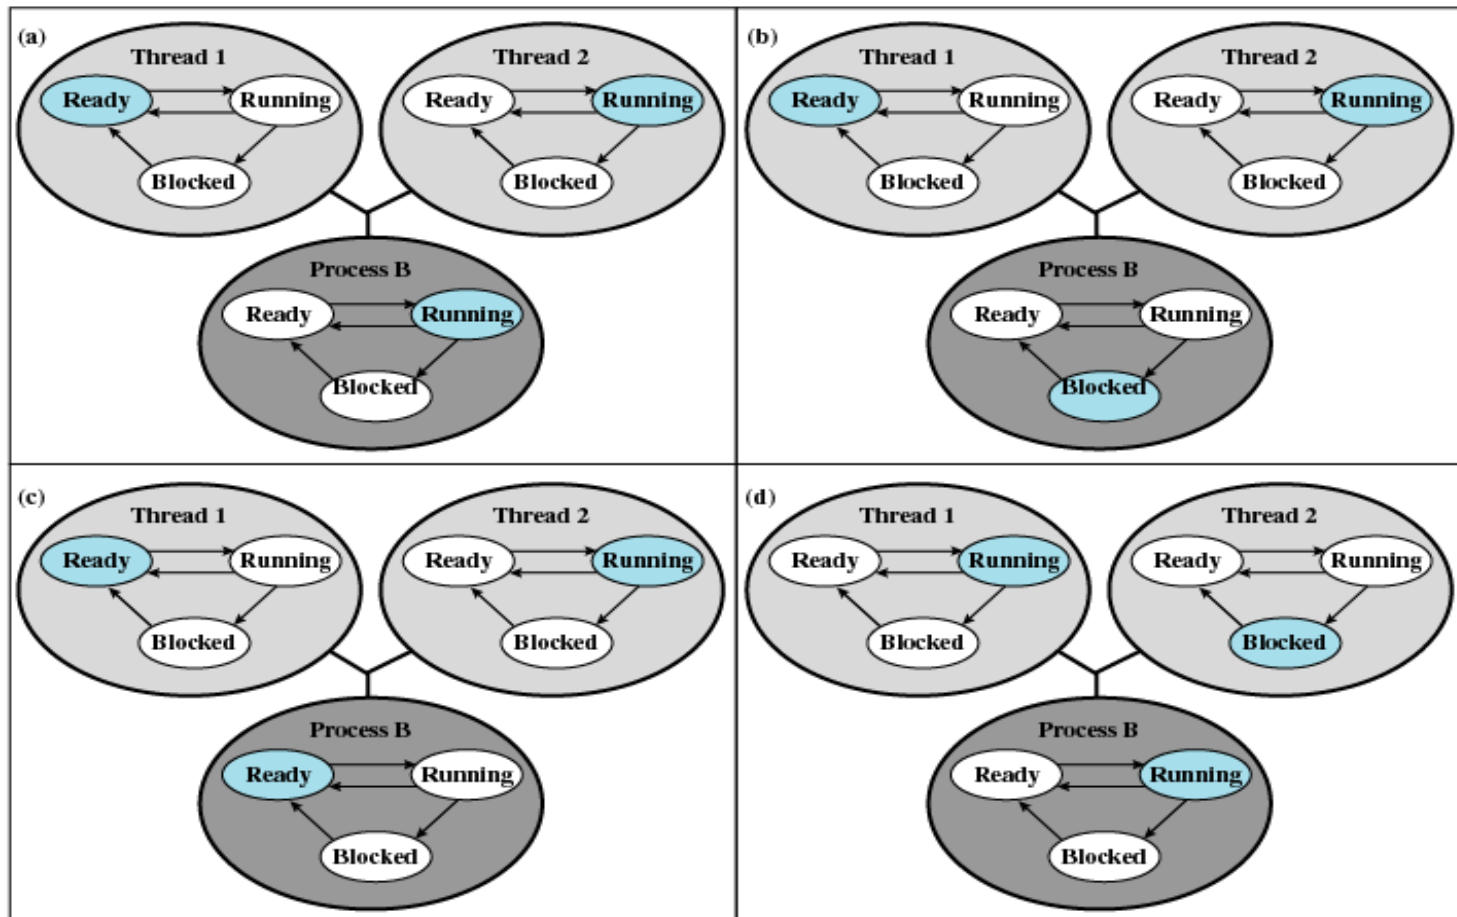
Multithreading



User-Level Threads

- All thread management is done by the application
- The kernel is not aware of the existence of threads






Colored state
is current state

Figure 4.7 Examples of the Relationships Between User-Level Thread States and Process States


Kernel-Level Threads

- All contemporary OS support kernel-level threads (such as Windows XP, Linux, Mac OS)
- Kernel maintains context information for the process and the threads
- Scheduling is done on a thread basis



Combined Approaches

- Example is Solaris
- Thread creation done in the user space
- Bulk of scheduling and synchronization of threads within application




Multithreading Models


(Relationship between KLT and ULT)

■ Many-to-One Model

- Many user-level threads mapped to single kernel thread




■ One-to-One Model



Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads



Thread Cancellation

- Terminating a thread before it has finished
 - Ex: User presses a stop button on a web browser that stop a web page from loading any further
- Cancellation of a target thread may occur in two different scenarios:
 - **Asynchronous cancellation.** One thread terminates the target thread immediately
 - **Deferred cancellation** allows the target thread to periodically check if it should be cancelled

Thread Pools

- Create a number of threads in a pool where they await work
- Advantages:
 - Usually slightly faster to service a request with an existing thread than create a new thread
 - Allows the number of threads in the application(s) to be bound to the size of the pool

Thread Libraries

- **Thread library** provides programmer with API for creating and managing threads
- Two primary ways of implementing
 - Library entirely in user space
 - Kernel-level library supported by the OS

Pthreads

- May be provided either as user-level or kernel-level
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (Solaris, Linux, Mac OS X)

Pthreads Example

```
#include <pthread.h>
#include <stdio.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* the thread */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    if (argc != 2) {
        fprintf(stderr, "usage: a.out <integer value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr, "%d must be >= 0\n", atoi(argv[1]));
        return -1;
    }
}
```


Pthreads Example (Cont.)

```
/* get the default attributes */
pthread_attr_init(&attr);
/* create the thread */
pthread_create(&tid,&attr,runner,argv[1]);
/* wait for the thread to exit */
pthread_join(tid,NULL);

printf("sum = %d\n",sum);
}

/* The thread will begin control in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```

Figure 4.9 Multithreaded C program using the Pthreads API.

Signal Handling

- Used in UNIX systems to notify a process that a particular event has occurred
- **Synchronuous Signals** : delivered to the same process that performed the operation
 - illegal memory access, division by zero
- **Asynchronuos Signals**: generated by an event external to a running process
 - terminating a process with specific keystrokes (such as <control><C>)

Signal Handling (2)

- All signals follow the following pattern:
 1. Signal is generated by particular event
 2. Signal is delivered to a process
 3. Once delivered, the signal must be handled.

- Handling Signals in multithreaded programs
 - Deliver the signal to the thread to which the signal applies
 - Deliver the signal to every thread in the process
 - Deliver the signal to certain threads in the process
 - Assign a specific thread to receive all signals for the process