

# CSE333 LAB

## Processes in Unix

Sanem Arslan Yılmaz

Zuhal Altuntaş

[Adopted from UNIX SYSTEMS Programming: Communication, Concurrency, and Threads]

# Process Identification

- A process is an instance of a program that is executing.
- UNIX identifies processes by a unique integer value called the process ID (pid)
  - Each process also has a parent process ID, which is initially the process ID of the process that created it.
- The **getpid()** and **getppid()** functions return the process ID and parent process ID

```
#include <unistd.h>
```

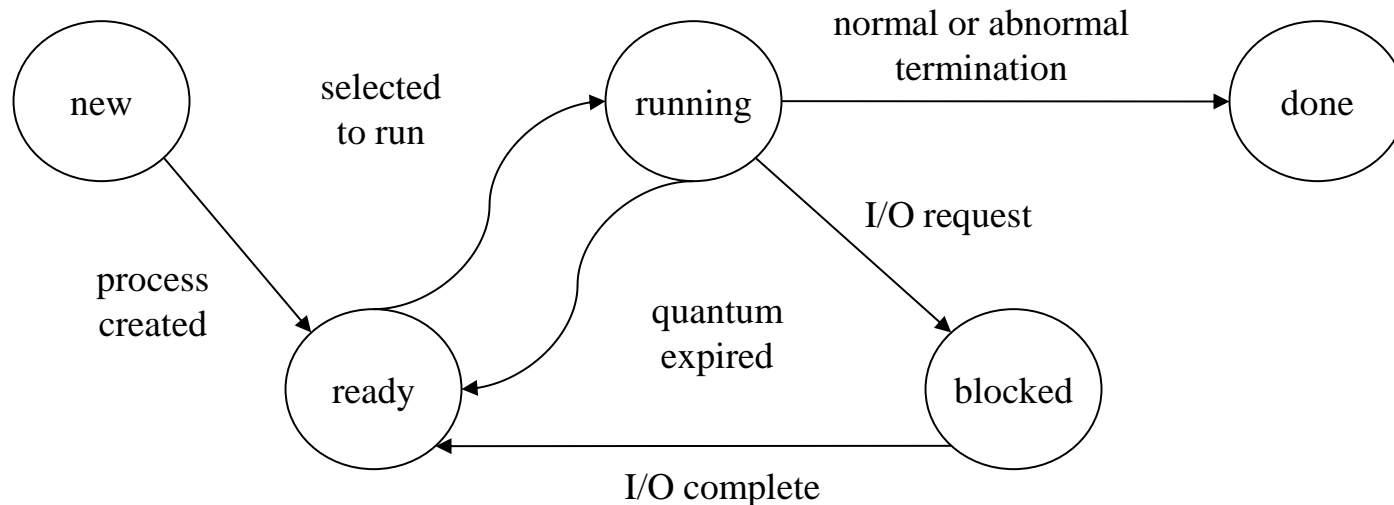
```
pid_t getpid(void);  
pid_t getppid(void);
```

- Example use:

```
printf("My process ID is %ld\n", getpid());
```

# Process State

- The state of a process indicates its status at a particular time
- Most operating systems use the following process states:
  - **New**: being created
  - **Running**: instructions are being executed
  - **Blocked**: waiting for an event such as I/O
  - **Ready**: waiting to be assigned to a processor
  - **Done**: finished execution (terminated)



# ps Utility

- The **ps** (**process status**) utility displays information about processes currently handled by the OS
- The **ps** utility is executed at a UNIX shell prompt; by default it displays information about processes for the current user.

- Column headings

UID	- user ID	STIME	- starting time of the process
PID	- process ID	TTY	- controlling terminal
PPID	- parent process ID	TIME	- cumulative execution time
C	- (obsolete)	CMD	- command name

- Example output produced by **ps -ef**

UID	PID	PPID	C	STIME	TTY	TIME	CMD
root	0	0	0	May 20	?	0:02	sched
root	1	0	0	May 20	?	2:38	/etc/init -
root	2	0	0	May 20	?	0:00	pageout
root	3	0	1	May 20	?	173:15	fsflush
root	433	1	0	May 20	console	0:00	/usr/bin/login
root	13259	436	0	13:54:49	?	0:00	/usr/lib/ssh/sshd
jjt107	13603	13261	0	13:55:29	pts/3	0:00	ps -ef
root	2017	2398	0	May 22	pts/5	0:00	sh
root	4210	5144	0	May 20	pts/1	0:00	sh
dan	28527	28500	0	May 22	pts/6	0:00	-sh
jjt107	13261	13259	0	13:54:59	pts/3	0:00	-csh
root	28499	436	0	May 22	?	0:00	/usr/lib/ssh/sshd
root	3110	436	0	May 25	?	0:00	/usr/lib/ssh/sshd
root	11090	436	0	May 23	?	0:00	/usr/lib/ssh/sshd

# UNIX Process Creation

- A process creates another process by calling the **fork()** function
  - The calling process is called the parent and the created process is called the child
  - The fork function copies the parent's memory image so that the new process receives a copy of the address space of the parent
  - Both processes continue at the instruction directly after the statement containing the `fork()` call (executing in their respective memory images)

```
#include <unistd.h>
```

```
pid_t fork(void);
```

- The return value from the **fork** () function is used to determine which process is the parent and which is the child; the child gets the value 0 while the parent gets the child's process ID
  - When the **fork** () call fails, it returns -1 and sets `errno` (a child is not created)

# Fork Example #1 (ch03/twoprocs.c)

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
```

```
int main(void) {
    pid_t childpid;

    childpid = fork();
    if (childpid == -1) {
        perror("Failed to fork");
        return 1;
    }
    if (childpid == 0)                /* child code */
        printf("I am child %ld\n", (long)getpid());
    else                             /* parent code */
        printf("I am parent %ld\n", (long)getpid());
    return 0;
}
```

## How to Compile:

```
gcc twoprocs.c -o twoprocs.out
```

## How to Run:

```
./twoprocs.out
```

# Fork Example #2 (ch03/badprocessID.c)

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
```

```
int main(void) {
    pid_t childpid;
    pid_t mypid;

    mypid = getpid();
    childpid = fork();
    if (childpid == -1) {
        perror("Failed to fork");
        return 1;
    }
    if (childpid == 0)                /* child code */
        printf("I am child %ld, ID = %ld\n", (long)getpid(), (long)mypid);
    else                             /* parent code */
        printf("I am parent %ld, ID = %ld\n", (long)getpid(), (long)mypid);
    return 0;
}
```

## How to Compile:

```
gcc badprocessID.c -o badprocessID.out
```

## How to Run:

```
./badprocessID.out
```

# Fork Example #3 (ch03/simplechain.c)

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
```

```
int main (int argc, char *argv[]) {
    pid_t childpid = 0;
    int i, n;
```

```
    if (argc != 2){ /* check for valid number of command-line arguments */
        fprintf(stderr, "Usage: %s processes\n", argv[0]);
        return 1;
    }
    n = atoi(argv[1]);
    for (i = 1; i < n; i++)
        if (childpid = fork())
            break;
```

```
    fprintf(stderr, "i:%d process ID:%ld parent ID:%ld child ID:%ld\n", i, (long)getpid(), (long)getppid(), (long)childpid);
    return 0;
}
```

## How to Compile:

```
gcc simplechain.c -o simplechain.out
```

## How to Run:


```
./simplechain.out 4
```



# Fork Example #3 (ch03/simplechain.c)

## Sample Run

```
% ./simplechain.out 4  
i:1 process ID: 2736 parent ID: 40 child ID: 3488  
i:2 process ID: 3488 parent ID: 2736 child ID: 512  
i:4 process ID: 120 parent ID: 512 child ID: 0  
i:3 process ID: 512 parent ID: 3488 child ID: 120
```



# Fork Example #4 (ch03/simplefan.c)

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
```

```
int main (int argc, char *argv[]) {
    pid_t childpid = 0;
    int i, n;
```

```
    if (argc != 2){ /* check for valid number of command-line arguments */
        fprintf(stderr, "Usage: %s processes\n", argv[0]);
        return 1;
    }
    n = atoi(argv[1]);
    for (i = 1; i < n; i++)
        if ((childpid = fork()) <= 0)
            break;
```

```
    fprintf(stderr, "i:%d process ID:%ld parent ID:%ld child ID:%ld\n", i, (long)getpid(), (long)getppid(), (long)childpid);
    return 0;
}
```

## How to Compile:

```
gcc simplefan.c -o simplefan.out
```


## How to Run:

```
./simplefan.out 4
```

# Fork Example #3 (ch03/simplefan.c)

## Sample Run

```
% ./simplefan.out 4  
i:1 process ID: 2736 parent ID: 120 child ID: 0  
i:2 process ID: 3488 parent ID: 120 child ID: 0  
i:4 process ID: 120 parent ID: 40 child ID: 512  
i:3 process ID: 512 parent ID: 120 child ID: 0
```




# Exercise 1

- Trace the following program segment and determine how many processes are created.

```
c2 = 0;  
c1 = fork(); /* fork number 1 */  
if (c1 == 0)  
c2 = fork(); /* fork number 2 */  
fork(); /* fork number 3 */  
if (c2 > 0)  
fork(); /* fork number 4 */
```

## Answer



# The wait() Function

- When a process creates a child, both parent and child proceed with execution from the point of the fork
- The parent process can run the **wait()** or **waitpid()** functions to block its execution until the child process finishes
- The **wait()** function causes the caller (i.e., the parent) to suspend execution until a child's status becomes available
  - A process status most commonly becomes available after process termination

```
#include <sys/wait.h>
```

```
pid_t wait(int *status_location);
```

- It takes one parameter, a pointer to the location for returning the status of the process
- The function returns either the process ID that terminated or `-1` (and sets `errno`)
- If a parent terminates without waiting for a child, the child becomes an **orphan** and is adopted by a special system process
- If a child process terminates and its parent does not wait for it, it becomes a **zombie** in UNIX terminology
  - Zombies stay in the system until they are waited for
- Traditionally, this process is called **init** and has process ID of 1; it periodically waits for children, so eventually orphaned zombies are removed

# Wait Example #1 (ch03/parentwaitpid.c)

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/wait.h>
```

```
int main (void) {
    pid_t childpid;

    childpid = fork();
    if (childpid == -1) {
        perror("Failed to fork");
        return 1;
    }
    if (childpid == 0)
        fprintf(stderr, "I am child %ld\n", (long)getpid());
    else if (wait(NULL) != childpid)
        fprintf(stderr, "A signal must have interrupted the wait!\n");
    else
        fprintf(stderr, "I am parent %ld with child %ld\n", (long)getpid(),
            (long)childpid);
    return 0;
}
```

## How to Compile:

```
gcc parentwaitpid.c -o parentwaitpid.out
```

## How to Run:

```
./parentwaitpid.out
```

# Wait Example #2 (ch03/fanwait.c)

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>
#include "restart.h"
```

```
int main(int argc, char *argv[]) {
    pid_t childpid;
    int i, n;

    if (argc != 2) {
        fprintf(stderr, "Usage: %s n\n", argv[0]);
        return 1;
    }
    n = atoi(argv[1]);
    for (i = 1; i < n; i++)
        if ((childpid = fork()) <= 0)
            break;

    while(wait(NULL) > 0) ; /* wait for all of your children */
    fprintf(stderr, "i:%d process ID:%ld parent ID:%ld child ID:%ld\n",
        i, (long)getpid(), (long)getppid(), (long)childpid);
    return 0;
}
```

## How to Compile:

```
gcc fanwait.c -o fanwait.out
```


## How to Run:

```
./fanwait.out 4
```

# wait() Example (using fan code)

## Sample Run

```
% ./fanwait.out 4
i:1 process ID: 2736 parent ID: 120 child ID: 0
i:2 process ID: 3488 parent ID: 120 child ID: 0
i:3 process ID: 512 parent ID: 120 child ID: 0
i:4 process ID: 120 parent ID: 40 child ID: 512
```





# The waitpid() Function

- The **waitpid()** function allows a parent to wait for a *particular* child to terminate

- It also allows a parent process to check whether a child has terminated without blocking

```
#include <sys/wait.h>
```

```
pid_t waitpid ( pid_t pid,  int *status_location,  int options );
```

- The function takes three parameters: a pid, a pointer to the location for returning a status, and a flag specifying options
- There are several variations on the **pid** parameter and the resulting actions of the **waitpid()** function
  - `pid > 0`        waits for the specific child whose process ID is pid
  - `pid == -1`    waits for any child
- By default, **waitpid()** waits only for terminated children, but this behavior is modifiable via the *options* argument, as described below.
- The **options** parameter is the bitwise inclusive OR of one or more flags
  - WNOHANG option causes the function to return even if the status of a child is not immediately available
  - It returns 0 to report that there are possibly unwaited-for children but that their status is not available

# waitpid() Example

- The following code segment waits for all children that have finished but avoids blocking if there are no children whose status is available

```
pid_t childpid;

// Wait for any child
childpid = waitpid(-1, NULL, WNOHANG);
while (childpid != 0)
{
    if ((childpid == -1) && (errno != EINTR))
        break;
    else
        childpid = waitpid(-1, NULL, WNOHANG);
} // End while
```


# Exercise 2

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main(void) {
    pid_t childpid;

    childpid = fork();
    if (childpid == -1) {
        perror("Failed to fork");
        return 1;
    }
    if (childpid == 0) {
        printf("A\n");
        printf("B\n");
    }
    else {
        printf("C\n");
        wait(NULL);
        printf("D\n");
    }
    return 0;
}
```

Output:



# The exec Family of Functions

- The **fork** ( ) function creates a copy of the calling process, but many applications require the child process to execute code that is different than that of the parent
- The **exec** family of functions provides a facility of overlaying the process image of the calling process with a new image
  - Usually the parent continues running the same code after the **fork** ( ) call, while the child process runs the new program (by means of an exec function call)
- There are six variations of the **exec** function
  - Each differ in the way command-line arguments and the environment are passed
  - They also differ in whether a full pathname must be given for the executable
- All exec functions returns **-1** if unsuccessful
  - If any of the exec functions return at all, the call was unsuccessful
- The **exec1**, **exec1p**, and **execle** functions pass the command-line arguments in an explicit list and are useful if the programmer knows the number of command line arguments at compile time
- The **execv**, **execvp**, and **execve** functions pass the command-line arguments in an argument array

# Exec Functions

```
#include <unistd.h>
```

```
extern char **environ;
```

```
int execl (const char *path, const char *arg0, ... /*, char *(0) */);
```

```
int execl (const char *path, const char *arg0, ... /*, char *(0), char *const envp[] */);
```

```
int execlp (const char *file, const char *arg0, ... /*, char *(0) */);
```

```
int execv(const char *path, char *const argv[]);
```

```
int execve (const char *path, char *const argv[], char *const envp[]);
```

```
int execvp (const char *file, char *const argv[]);
```

# execl() Example (ch03/execls.c)

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

int main (void)
{
    pid_t childpid;

    childpid = fork();
    if (childpid == -1)
    {
        perror("Fork failed");
        return 1;
    } // End if

    if (childpid == 0)
    {
        execl("/bin/ls", "ls", "-l", NULL);
        perror("Child failed when running execl");
        return 1;
    } // End if

    if (childpid != wait(NULL))
    {
        perror("Parent failed to wait due to signal or error");
        return 1;
    } // End if

    return 0;
} // End main
```

This program creates a child process that runs the ls utility using the "-l" option. It displays a long listing of the contents of the current working directory

• "/bin/ls"	const char *path
• "ls"	const char *arg0
• "-l"	const char *arg1
• NULL	NULL pointer

## How to Compile:

gcc execls.c -o execls.out

## How to Run:

./execls.out

# execvp() Example (ch03/execcmd.c)

```
int main (int argc, char *argv[])
{
    pid_t childpid;

    if (argc < 2)
    {
        fprintf(stderr, "Usage: a.out command arg1 arg2 ...\n");
        return 1;
    } // end if

    childpid = fork();
    if (childpid == -1)
    {
        perror("Fork failed");
        return 1;
    } // End if

    if (childpid == 0) /* Child code */
    {
        execvp(argv[1], &argv[1]);
        perror("Child failed upon running execvp function");
        return 1;
    }

    if (childpid != wait(NULL)) /* Parent code */
    {
        perror("Parent failed to wait due to signal or error");
        return 1;
    }

    return 0;
} // End main
```

This program creates a child process that runs the command or program submitted on the command line.

- |            |                    |
|------------|--------------------|
| • argv[1]  | const char *file   |
| • &argv[1] | const char *argv[] |

## How to Compile:

gcc execcmd.c -o execcmd.out

## How to Run:

./execcmd.out ls

./execcmd.out ls -l

# Background Processes

- A **shell** in UNIX terminology is a command interpreter that provides a prompt for a command, reads the command from standard input, forks a child to execute the command and waits for the child to finish
- When standard input and output come from a terminal type or device, a user can terminate an executing command by entering the interrupt character (commonly Ctrl-C)
- Most command shells interpret an input line ending with **&** (i.e., ampersand) as a command that should be executed as a background process
  - Example: `./a.out &`
  - When a shell creates a background process, it does not wait for the process to complete before issuing a prompt and accepting additional commands
  - Also, Ctrl-C from the keyboard does not terminate a background process
- A **daemon** is a background process that normally runs indefinitely
- UNIX relies on many daemon processes to perform routine tasks
  - For example, the Solaris `pageout` daemon handles paging for memory management
  - The `in.rlogind` daemon handles remote login requests
  - Other daemons handle mail, file transfer, statistics and printer requests

