

CSE333 LAB

Redirection and Pipes

Zuhal Altuntaş

[Slides by Dr. Sanem Arslan Yılmaz]

I/O Redirection

- Redirect standard output (create / truncate)
`a.out > outfile`
- Redirect standard output (create / append)
`a.out >> outfile`
- Redirect standard input
`a.out < inputfile`
- Redirect standard input and standard output
`a.out < inputfile > outputfile`

Standard File Descriptors

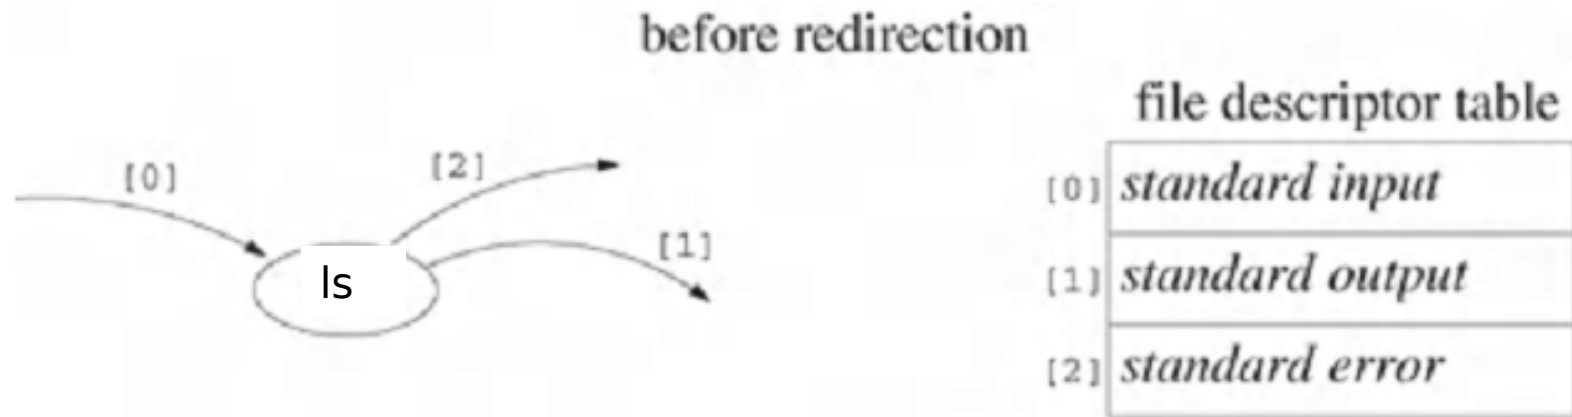
- When a program is executed, it automatically starts with **three** open file streams

Symbolic Name	I/O Name	Meaning	Default
• STDIN_FILENO	stdin	standard input device	keyboard
• STDOUT_FILENO	stdout	standard output device	monitor
• STDERR_FILENO	stderr	standard error report device	monitor
•			

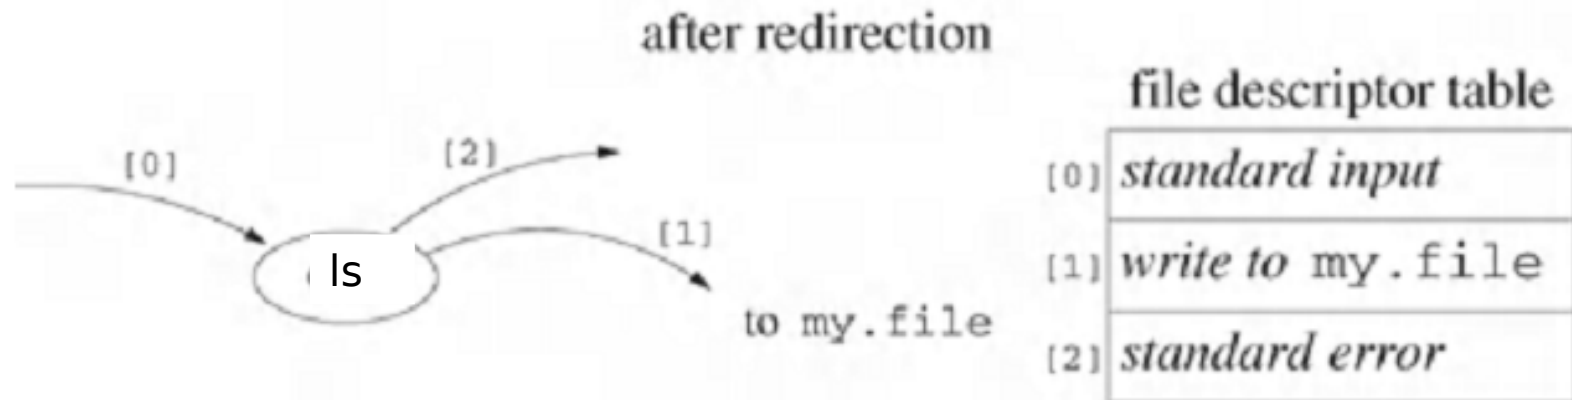
file descriptor table	
[0]	<i>standard input</i>
[1]	<i>standard output</i>
[2]	<i>standard error</i>

File Descriptor Table

\$ ls



\$ ls > myfile



dup2() function

SYNOPSIS

```
#include <unistd.h>
```

```
int dup2(int fildes, int fildes2);
```

- The **dup2** function takes two parameters, **fildes** and **fildes2**.
- It closes entry **fildes2** of the file descriptor table if it was open and then copies the pointer of entry **fildes** into entry **fildes2**.
- On success, **dup2** returns the file descriptor value that was duplicated. On failure, **dup2** returns -1.

Example (ch06/redirect.c)

```
#include <fcntl.h>
#include <stdio.h>
#include <sys/stat.h>
#include <unistd.h>
```

```
#define CREATE_FLAGS (O_WRONLY | O_CREAT | O_APPEND)
#define CREATE_MODE (S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH)
```

```
int main(void) {
    int fd;

    fd = open("my.file", CREATE_FLAGS, CREATE_MODE);
    if (fd == -1) {
        perror("Failed to open my.file");
        return 1;
    }
    if (dup2(fd, STDOUT_FILENO) == -1) {
        perror("Failed to redirect standard output");
        return 1;
    }
    if (close(fd) == -1) {
        perror("Failed to close the file");
        return 1;
    }
    printf("Output will be seen in my.file\n");
    return 0;
}
```

How to Compile:

```
gcc redirect.c -o redirect.out
```

How to Run:

```
./redirect.out
```

Status of file descriptor table

after open

file descriptor table

[0]	<i>standard input</i>
[1]	<i>standard output</i>
[2]	<i>standard error</i>
[3]	<i>write to my.file</i>

after dup2

file descriptor table


[0]	<i>standard input</i>
[1]	<i>write to my.file</i>
[2]	<i>standard error</i>
[3]	<i>write to my.file</i>

after close

file descriptor table

[0]	<i>standard input</i>
[1]	<i>write to my.file</i>
[2]	<i>standard error</i>

Pipelines

- The standard output of one *command* is fed into the standard input of another with *pipelines*.
 - ls -l | sort
 - 

Pipelines

- The standard output of one *command* is fed into the standard input of another with *pipelines*.
 - ls -l | sort → ls -l > file_list.txt
sort < file_list.txt

Pipelines

- The standard output of one *command* is fed into the standard input of another with *pipelines*.

- `ls -l | sort` → `ls -l > file_list.txt`
• `sort < file_list.txt`

- `ls -l | sort | wc -l`

pipe() Function

- The simplest mechanism in UNIX for interprocess communication is the unnamed pipe, which is represented by a special file
- The `pipe()` function creates a communication buffer that the caller can access through the two-entry array parameter (i.e., `fileDescriptors[2]`)
- The data written to `fileDescriptors[1]` can be then read from `fileDescriptors[0]` on a first-in-first-out basis
-
- `#include <unistd.h>`
-
- `int pipe(int fileIDs[2]);` // An array with two members
-
- If successful, the function returns zero; otherwise, it returns -1 and sets `errno`

pipe() Function (continued)

- A pipe has no external or permanent name, so a program can access it only through its two descriptors
 - For this reason, a pipe can be used only by the process that created it and by descendants that inherit the descriptors by way of a `fork()` call

pipe() Function (continued)

- A pipe has no external or permanent name, so a program can access it only through its two descriptors
 - For this reason, a pipe can be used only by the process that created it and by descendants that inherit the descriptors by way of a `fork()` call
- The `pipe()` function creates a unidirectional communication buffer

pipe() Function (continued)

- A pipe has no external or permanent name, so a program can access it only through its two descriptors
 - For this reason, a pipe can be used only by the process that created it and by descendants that inherit the descriptors by way of a `fork()` call
- The `pipe()` function creates a unidirectional communication buffer
- When a process calls `read()` on a pipe, the `read()` function returns immediately if the pipe is not empty

pipe() Function (continued)

- A pipe has no external or permanent name, so a program can access it only through its two descriptors
 - For this reason, a pipe can be used only by the process that created it and by descendants that inherit the descriptors by way of a `fork()` call
- The `pipe()` function creates a unidirectional communication buffer
- When a process calls `read()` on a pipe, the `read()` function returns immediately if the pipe is not empty
- If the pipe is empty, the `read()` function blocks until something is written to the pipe, as long as some process has the pipe open for writing

pipe() Function (continued)

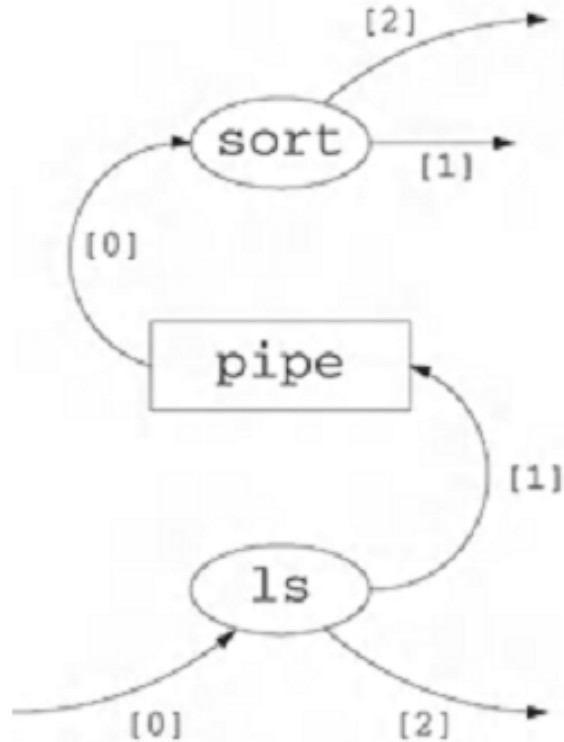
- A pipe has no external or permanent name, so a program can access it only through its two descriptors
 - For this reason, a pipe can be used only by the process that created it and by descendants that inherit the descriptors by way of a `fork()` call
- The `pipe()` function creates a unidirectional communication buffer
- When a process calls `read()` on a pipe, the `read()` function returns immediately if the pipe is not empty
- If the pipe is empty, the `read()` function blocks until something is written to the pipe, as long as some process has the pipe open for writing
- On the other hand, if no process has the pipe open for writing, a `read()` call on an empty pipe returns zero, indicating end of file

pipe() Function (continued)

- A pipe has no external or permanent name, so a program can access it only through its two descriptors
 - For this reason, a pipe can be used only by the process that created it and by descendants that inherit the descriptors by way of a `fork()` call
- The `pipe()` function creates a unidirectional communication buffer
- When a process calls `read()` on a pipe, the `read()` function returns immediately if the pipe is not empty
- If the pipe is empty, the `read()` function blocks until something is written to the pipe, as long as some process has the pipe open for writing
- On the other hand, if no process has the pipe open for writing, a `read()` call on an empty pipe returns zero, indicating end of file
- Normally, a parent process uses one or more pipes to communicate with its children as shown on the next slide
 - A parent creates a pipe before calling `fork()` to create a child
 - The parent then writes a message to the pipe
 - The child reads a message from the pipe

Status of file descriptor table

```
$ ls -l | sort
```



sort

file descriptor table

[0]	pipe read
[1]	standard output
[2]	standard error

ls

file descriptor table

[0]	standard input
[1]	pipe write
[2]	standard error

Example (ch06/simpleredirect.c)

```
#include <errno.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>

int main(void) {
    pid_t childpid;
    int fd[2];
    if ((pipe(fd) == -1) || ((childpid = fork()) == -1)) {
        perror("Failed to setup pipeline");
        return 1;
    }
    if (childpid == 0) { /* ls is the child */
        if (dup2(fd[1], STDOUT_FILENO) == -1)
            perror("Failed to redirect stdout of ls");
        else if ((close(fd[0]) == -1) || (close(fd[1]) == -1))
            perror("Failed to close extra pipe descriptors on ls");
        else {
            execl("/bin/ls", "ls", "-l", NULL);
            perror("Failed to exec ls");
        }
        return 1;
    }
    if (dup2(fd[0], STDIN_FILENO) == -1) /* sort is the parent */
        perror("Failed to redirect stdin of sort");
    else if ((close(fd[0]) == -1) || (close(fd[1]) == -1))
        perror("Failed to close extra pipe file descriptors on sort");
    else {
        execl("/bin/sort", "sort", "-n", "+4", NULL);
        perror("Failed to exec sort");
    }
    return 1;
}
```

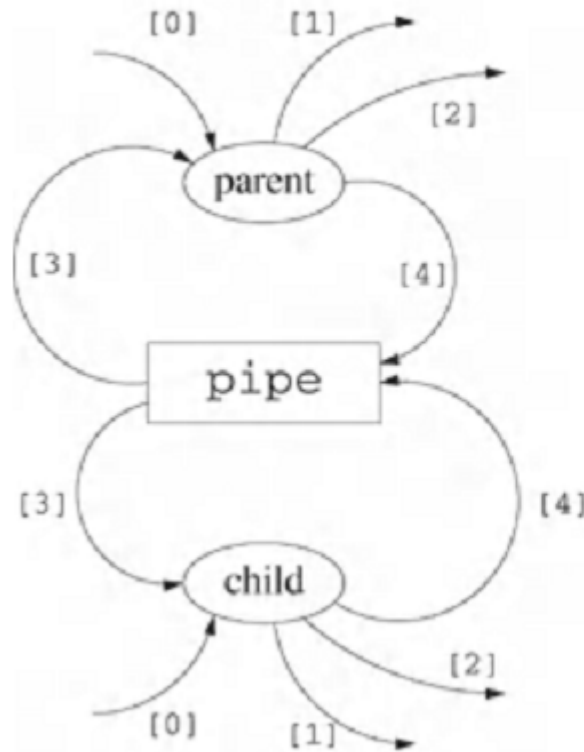
How to Compile:

```
gcc simpleredirect.c -o
simpleredirect.out
```

How to Run:

```
./simpleredirect.out
```

Status of file descriptor table (after fork() function)



parent

file descriptor table

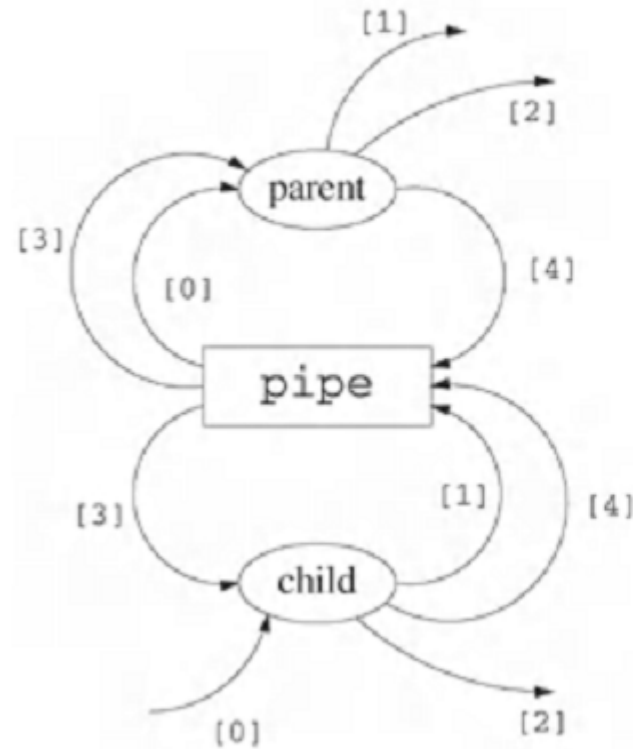
[0]	<i>standard input</i>
[1]	<i>standard output</i>
[2]	<i>standard error</i>
[3]	<i>pipe read</i>
[4]	<i>pipe write</i>

child

file descriptor table

[0]	<i>standard input</i>
[1]	<i>standard output</i>
[2]	<i>standard error</i>
[3]	<i>pipe read</i>
[4]	<i>pipe write</i>

Status of file descriptor table (after both dup2() functions)



parent

file descriptor table

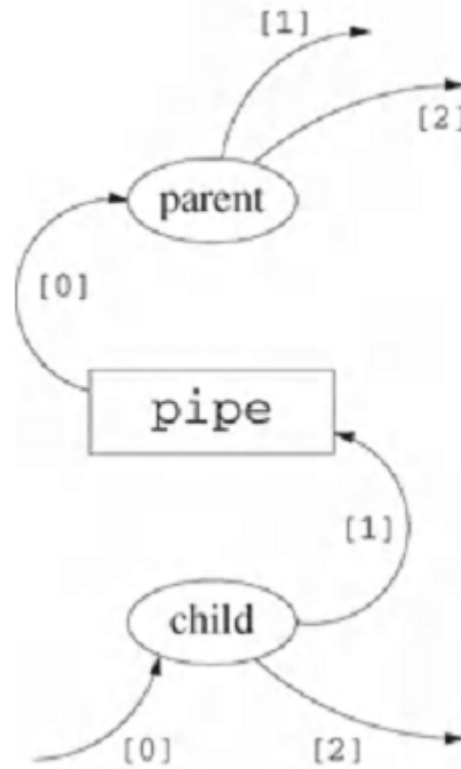
[0]	pipe read
[1]	standard output
[2]	standard error
[3]	pipe read
[4]	pipe write

child

file descriptor table

[0]	standard input
[1]	pipe write
[2]	standard error
[3]	pipe read
[4]	pipe write

Status of file descriptor table (after all close() functions)



parent

file descriptor table

[0]	<i>pipe read</i>
[1]	<i>standard output</i>
[2]	<i>standard error</i>

child

file descriptor table

[0]	<i>standard input</i>
[1]	<i>pipe write</i>
[2]	<i>standard error</i>